

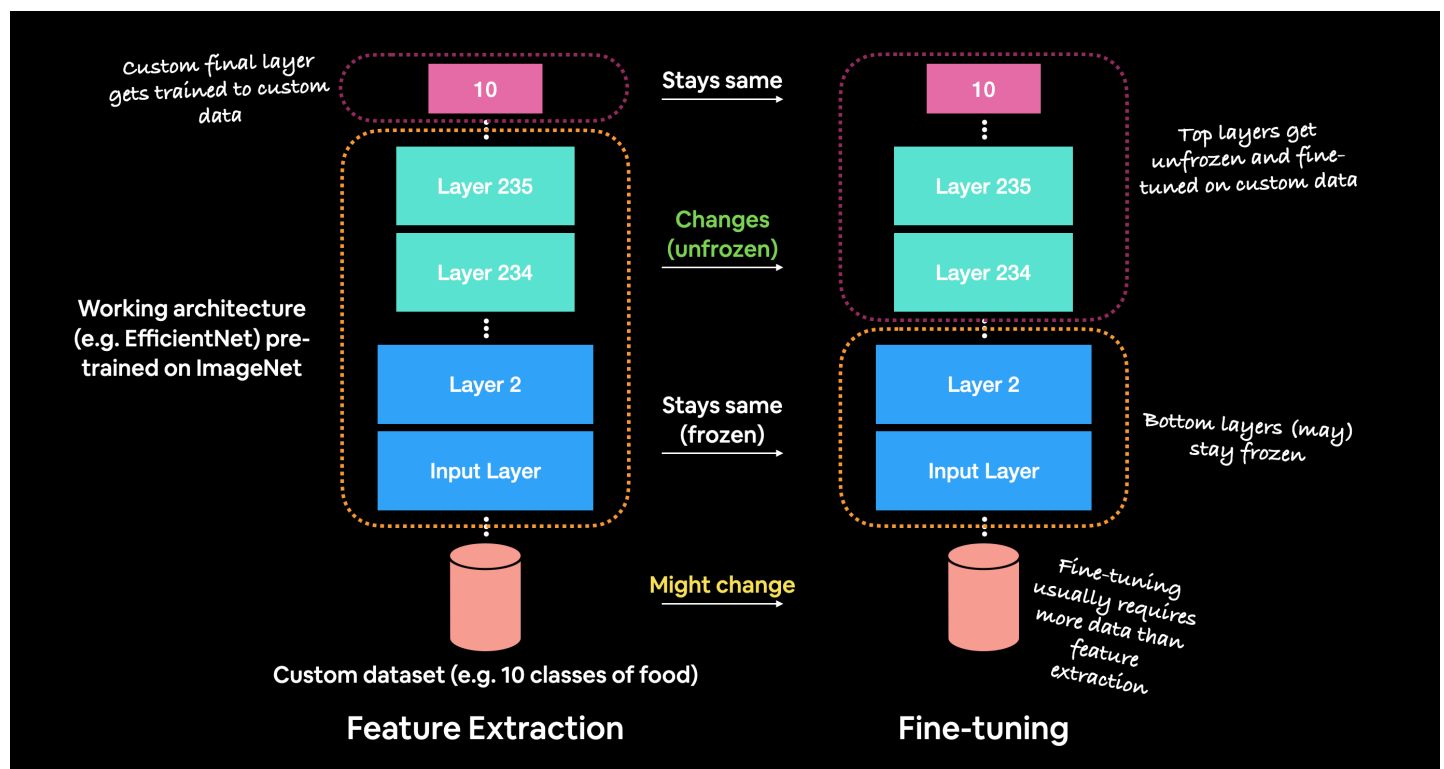
## 05. Transfer Learning with TensorFlow Part 2: Fine-tuning

In the previous section, we saw how we could leverage feature extraction transfer learning to get far better results on our Food Vision project than building our own models (even with less data).

Now we're going to cover another type of transfer learning: fine-tuning.

In **fine-tuning transfer learning** the pre-trained model weights from another model are unfrozen and tweaked during to better suit your own data.

For feature extraction transfer learning, you may only train the top 1-3 layers of a pre-trained model with your own data, in fine-tuning transfer learning, you might train 1-3+ layers of a pre-trained model (where the '+' indicates that many or all of the layers could be trained).



*Feature extraction transfer learning vs. fine-tuning transfer learning. The main difference between the two is that in fine-tuning, more layers of the pre-trained model get unfrozen and tuned on custom data. This fine-tuning usually takes more data than feature extraction to be effective.*

### What we're going to cover

We're going to go through the follow with TensorFlow:

- Introduce fine-tuning, a type of transfer learning to modify a pre-trained model to be more suited to your data
- Using the Keras Functional API (a differnt way to build models in Keras)
- Using a smaller dataset to experiment faster (e.g. 1-10% of training samples of 10 classes of food)
- Data augmentation (how to make your training dataset more diverse without adding more data)
- Running a series of modelling experiments on our Food Vision data
  - Model 0: a transfer learning model using the Keras Functional API
  - Model 1: a feature extraction transfer learning model on 1% of the data with data augmentation
  - Model 2: a feature extraction transfer learning model on 10% of the data with data augmentation
  - Model 3: a fine-tuned transfer learning model on 10% of the data
  - Model 4: a fine-tuned transfer learning model on 100% of the data
- Introduce the ModelCheckpoint callback to save intermediate training results
- Compare model experiments results using TensorBoard

# How you can use this notebook

You can read through the descriptions and the code (it should all run, except for the cells which error on purpose), but there's a better option.

Write all of the code yourself.

Yes. I'm serious. Create a new notebook, and rewrite each line by yourself. Investigate it, see if you can break it, why does it break?

You don't have to write the text descriptions but writing the code yourself is a great way to get hands-on experience.

Don't worry if you make mistakes, we all do. The way to get better and make less mistakes is to **write more code**.

In [ ]:

```
# Are we using a GPU? (if not & you're using Google Colab, go to Runtime -> Change Runtime Type -> Hardware Accelerator: GPU )
!nvidia-smi
```

Tue Feb 16 02:14:29 2021

```
+-----+
| NVIDIA-SMI 460.39          Driver Version: 460.32.03          CUDA Version: 11.2          |
+-----+-----+-----+-----+-----+-----+
| GPU   Name                Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan   Temp   Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|                                           MIG M.           |
+-----+-----+-----+-----+-----+-----+
|    0   Tesla T4              Off      | 00000000:00:04.0 Off |                    0 |
| N/A    73C    P8             13W /  70W |  0MiB / 15109MiB |      0%      Default |
|                                           | N/A              |
+-----+-----+-----+-----+-----+-----+

```

```
+-----+
| Processes:
| GPU   GI    CI          PID    Type    Process name                        GPU Memory
|      ID    ID                                   |          Usage
+-----+-----+-----+-----+-----+-----+
|   No running processes found
+-----+

```

## Creating helper functions

Throughout your machine learning experiments, you'll likely come across snippets of code you want to use over and over again.

For example, a plotting function which plots a model's `history` object (see `plot_loss_curves()` below).

You could recreate these functions over and over again.

But as you might've guessed, rewriting the same functions becomes tedious.

One of the solutions is to store them in a helper script such as `helper_functions.py`. And then import the necessary functionality when you need it.

For example, you might write:

```
from helper_functions import plot_loss_curves

...

plot_loss_curves(history)
```

Let's see what this looks like.

In [ ]:

```
# Get helper_functions.py script from course GitHub
!wget https://raw.githubusercontent.com/mrdbourke/tensorflow-deep-learning/main/extras/helper_functions.py

# Import helper functions we're going to use
from helper_functions import create_tensorboard_callback, plot_loss_curves, unzip_data, walk_through_dir

--2021-02-16 02:14:32-- https://raw.githubusercontent.com/mrdbourke/tensorflow-deep-learning/main/extras/helper_functions.py
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.111.133, 185.199.108.133, 185.199.109.133, ...
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.111.133|:443.. connected.
HTTP request sent, awaiting response... 200 OK
Length: 9373 (9.2K) [text/plain]
Saving to: 'helper_functions.py.1'

helper_functions.py 100%[=====>] 9.15K --.-KB/s in 0s

2021-02-16 02:14:32 (99.6 MB/s) - 'helper_functions.py.1' saved [9373/9373]
```

Wonderful, now we've got a bunch of helper functions we can use throughout the notebook without having to rewrite them from scratch each time.

**Note:** If you're running this notebook in Google Colab, when it times out Colab will delete the `helper_functions.py` file. So to use the functions imported above, you'll have to rerun the cell.

## 10 Food Classes: Working with less data

We saw in the [previous notebook](#) that we could get great results with only 10% of the training data using transfer learning with TensorFlow Hub.

In this notebook, we're going to continue to work with smaller subsets of the data, except this time we'll have a look at how we can use the in-built pretrained models within the `tf.keras.applications` module as well as how to fine-tune them to our own custom dataset.

We'll also practice using a new but similar dataloader function to what we've used before, `image_dataset_from_directory()` which is part of the `tf.keras.preprocessing` module.

Finally, we'll also be practicing using the [Keras Functional API](#) for building deep learning models. The Functional API is a more flexible way to create models than the `tf.keras.Sequential` API.

We'll explore each of these in more detail as we go.

Let's start by downloading some data.

In [ ]:

```
# Get 10% of the data of the 10 classes
!wget https://storage.googleapis.com/ztm_tf_course/food_vision/10_food_classes_10_percent.zip

unzip_data("10_food_classes_10_percent.zip")

--2021-02-16 02:14:53-- https://storage.googleapis.com/ztm_tf_course/food_vision/10_food_classes_10_percent.zip
Resolving storage.googleapis.com (storage.googleapis.com)... 172.217.12.240, 172.217.15.112, 172.253.62.128, ...
Connecting to storage.googleapis.com (storage.googleapis.com)|172.217.12.240|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 16954612 (16.1M) [application/zip]
```

```
Length: 168546183 (161M) [application/zip]
```

```
Saving to: '10_food_classes_10_percent.zip.1'
```

```
10_food_classes_10_ 100%[=====>] 160.74M 186MB/s in 0.9s
```

```
2021-02-16 02:14:54 (186 MB/s) - '10_food_classes_10_percent.zip.1' saved [168546183/168546183]
```

The dataset we're downloading is the 10 food classes dataset (from Food 101) with 10% of the training images we used in the previous notebook.

❏ **Note:** You can see how this dataset was created in the [image data modification notebook](#).

```
In [ ]:
```

```
# Walk through 10 percent data directory and list number of files
walk_through_dir("10_food_classes_10_percent")
```

```
There are 2 directories and 0 images in '10_food_classes_10_percent'.
There are 10 directories and 0 images in '10_food_classes_10_percent/train'.
There are 0 directories and 75 images in '10_food_classes_10_percent/train/ice_cream'.
There are 0 directories and 75 images in '10_food_classes_10_percent/train/ramen'.
There are 0 directories and 75 images in '10_food_classes_10_percent/train/chicken_wings'
.
There are 0 directories and 75 images in '10_food_classes_10_percent/train/pizza'.
There are 0 directories and 75 images in '10_food_classes_10_percent/train/steak'.
There are 0 directories and 75 images in '10_food_classes_10_percent/train/fried_rice'.
There are 0 directories and 75 images in '10_food_classes_10_percent/train/hamburger'.
There are 0 directories and 75 images in '10_food_classes_10_percent/train/grilled_salmon'
'.
There are 0 directories and 75 images in '10_food_classes_10_percent/train/sushi'.
There are 0 directories and 75 images in '10_food_classes_10_percent/train/chicken_curry'
.
There are 10 directories and 0 images in '10_food_classes_10_percent/test'.
There are 0 directories and 250 images in '10_food_classes_10_percent/test/ice_cream'.
There are 0 directories and 250 images in '10_food_classes_10_percent/test/ramen'.
There are 0 directories and 250 images in '10_food_classes_10_percent/test/chicken_wings'
.
There are 0 directories and 250 images in '10_food_classes_10_percent/test/pizza'.
There are 0 directories and 250 images in '10_food_classes_10_percent/test/steak'.
There are 0 directories and 250 images in '10_food_classes_10_percent/test/fried_rice'.
There are 0 directories and 250 images in '10_food_classes_10_percent/test/hamburger'.
There are 0 directories and 250 images in '10_food_classes_10_percent/test/grilled_salmon'
'.
There are 0 directories and 250 images in '10_food_classes_10_percent/test/sushi'.
There are 0 directories and 250 images in '10_food_classes_10_percent/test/chicken_curry'
.
.
```

We can see that each of the training directories contain 75 images and each of the testing directories contain 250 images.

Let's define our training and test filepaths.

```
In [ ]:
```

```
# Create training and test directories
train_dir = "10_food_classes_10_percent/train/"
test_dir = "10_food_classes_10_percent/test/"
```

Now we've got some image data, we need a way of loading it into a TensorFlow compatible format.

Previously, we've used the `ImageDataGenerator` class. And while this works well and is still very commonly used, this time we're going to use the `image_data_from_directory` function.

It works much the same way as `ImageDataGenerator`'s `flow_from_directory` method meaning your images need to be in the following file format:

Example of file structure

```
10_food_classes_10_percent <- top level folder
├── train <- training images
│   ├── pizza
│   │   ├── 1008104.jpg
│   │   ├── 1638227.jpg
│   │   └── ...
│   └── steak
│       ├── 1000205.jpg
│       ├── 1647351.jpg
│       └── ...
└── test <- testing images
    ├── pizza
    │   ├── 1001116.jpg
    │   ├── 1507019.jpg
    │   └── ...
    └── steak
        ├── 100274.jpg
        ├── 1653815.jpg
        └── ...
```

One of the main benefits of using `tf.keras.preprocessing.image_dataset_from_directory()` rather than `ImageDataGenerator` is that it creates a `tf.data.Dataset` object rather than a generator. The main advantage of this is the `tf.data.Dataset` API is much more efficient (faster) than the `ImageDataGenerator` API which is paramount for larger datasets.

Let's see it in action.

In [ ]:

```
# Create data inputs
import tensorflow as tf
IMG_SIZE = (224, 224) # define image size
train_data_10_percent = tf.keras.preprocessing.image_dataset_from_directory(directory=train_dir,
                                                                              image_size=
                                                                              label_mode=
                                                                              batch_size=
                                                                              32) # batch_size is 32 by default, this is generally a good number
test_data_10_percent = tf.keras.preprocessing.image_dataset_from_directory(directory=test_dir,
                                                                              image_size=I
                                                                              label_mode="
                                                                              categorical")
```

Found 750 files belonging to 10 classes.  
Found 2500 files belonging to 10 classes.

Wonderful! Looks like our dataloaders have found the correct number of images for each dataset.

For now, the main parameters we're concerned about in the `image_dataset_from_directory()` function are:

- `directory` - the filepath of the target directory we're loading images in from.
- `image_size` - the target size of the images we're going to load in (height, width).
- `batch_size` - the batch size of the images we're going to load in. For example if the `batch_size` is 32 (the default), batches of 32 images and labels at a time will be passed to the model.

There are more we could play around with if we needed to [in the tf.keras.preprocessing documentation](#).

If we check the training data datatype we should see it as a `BatchDataset` with shapes relating to our data.

In [ ]:

```
# Check the training data datatype
train_data_10_percent
```

Out[ ]:

```
<BatchDataset shapes: ((None, 224, 224, 3), (None, 10)), types: (tf.float32, tf.float32)>
```

In the above output:

- `(None, 224, 224, 3)` refers to the tensor shape of our images where `None` is the batch size, `224` is the height (and width) and `3` is the color channels (red, green, blue).
- `(None, 10)` refers to the tensor shape of the labels where `None` is the batch size and `10` is the number of possible labels (the 10 different food classes).
- Both image tensors and labels are of the datatype `tf.float32`.

The `batch_size` is `None` due to it only being used during model training. You can think of `None` as a placeholder waiting to be filled with the `batch_size` parameter from `image_dataset_from_directory()`.

Another benefit of using the `tf.data.Dataset` API are the associated methods which come with it.

For example, if we want to find the name of the classes we were working with, we could use the `class_names` attribute.

In [ ]:

```
# Check out the class names of our dataset
train_data_10_percent.class_names
```

Out[ ]:

```
['chicken_curry',
 'chicken_wings',
 'fried_rice',
 'grilled_salmon',
 'hamburger',
 'ice_cream',
 'pizza',
 'ramen',
 'steak',
 'sushi']
```

Or if we wanted to see an example batch of data, we could use the `take()` method.

In [ ]:

```
# See an example batch of data
for images, labels in train_data_10_percent.take(1):
    print(images, labels)
```

```
tf.Tensor(
[[[1.00000000e+00 0.00000000e+00 3.10000000e+01]
 [1.00000000e+00 0.00000000e+00 3.10000000e+01]
 [1.00000000e+00 0.00000000e+00 3.10000000e+01]
 ...
 [1.00000000e+00 0.00000000e+00 3.10000000e+01]
 [1.00000000e+00 0.00000000e+00 3.10000000e+01]
 [1.00000000e+00 0.00000000e+00 3.10000000e+01]]

[[1.00000000e+00 0.00000000e+00 3.10000000e+01]
 [1.00000000e+00 0.00000000e+00 3.10000000e+01]
 [1.00000000e+00 0.00000000e+00 3.10000000e+01]
 ...
 [1.00000000e+00 0.00000000e+00 3.10000000e+01]
 [1.00000000e+00 0.00000000e+00 3.10000000e+01]]
```

```
[1.00000000e+00 0.00000000e+00 3.10000000e+01]]

[[1.00000000e+00 0.00000000e+00 3.10000000e+01]
 [1.00000000e+00 0.00000000e+00 3.10000000e+01]
 [1.00000000e+00 0.00000000e+00 3.10000000e+01]
 ...
 [1.00000000e+00 0.00000000e+00 3.10000000e+01]
 [1.00000000e+00 0.00000000e+00 3.10000000e+01]
 [1.00000000e+00 0.00000000e+00 3.10000000e+01]]

...

[[1.07500107e+02 9.49286346e+01 8.90816803e+01]
 [1.15714394e+02 1.01285782e+02 9.41582489e+01]
 [1.17974548e+02 1.04188812e+02 9.51888123e+01]
 ...
 [1.18617378e+02 2.90000000e+01 2.76939182e+01]
 [1.19000000e+02 2.90000000e+01 2.80000000e+01]
 [1.20076530e+02 3.00765305e+01 2.98622665e+01]]

[[1.07045891e+02 9.20458908e+01 8.70458908e+01]
 [1.15852043e+02 1.00852043e+02 9.38520432e+01]
 [1.15841820e+02 1.02056099e+02 9.32703857e+01]
 ...
 [1.17943863e+02 2.99438648e+01 2.87296009e+01]
 [1.17923454e+02 2.79234543e+01 2.79234543e+01]
 [1.17857117e+02 2.78571167e+01 2.78571167e+01]]

[[1.00785606e+02 8.57856064e+01 8.07856064e+01]
 [1.13999931e+02 9.89999313e+01 9.28061218e+01]
 [1.11999931e+02 9.74284973e+01 9.06427841e+01]
 ...
 [1.16857086e+02 2.88570862e+01 2.76428223e+01]
 [1.14571381e+02 2.65713806e+01 2.55713806e+01]
 [1.14642822e+02 2.46428223e+01 2.66428223e+01]]]

[[[9.76530609e+01 7.90663223e+01 6.50663223e+01]
 [1.04392860e+02 8.82500000e+01 7.33214264e+01]
 [1.09520409e+02 9.69591827e+01 8.36683655e+01]
 ...
 [2.44423492e+02 2.37423492e+02 2.19423492e+02]
 [2.46760208e+02 2.39760208e+02 2.21760208e+02]
 [2.48000000e+02 2.41000000e+02 2.23000000e+02]]

[[1.06326530e+02 8.95204086e+01 7.75867386e+01]
 [1.11239799e+02 9.68877640e+01 8.40918427e+01]
 [1.50612259e+02 1.41198990e+02 1.26897972e+02]
 ...
 [2.40857162e+02 2.33857162e+02 2.15857162e+02]
 [2.45071442e+02 2.38071442e+02 2.20071442e+02]
 [2.47000000e+02 2.40000000e+02 2.22000000e+02]]

[[1.15163269e+02 1.02295921e+02 9.33724518e+01]
 [1.54887756e+02 1.45117355e+02 1.34989807e+02]
 [2.00071426e+02 1.92836746e+02 1.79790833e+02]
 ...
 [2.40357193e+02 2.33357193e+02 2.15357193e+02]
 [2.43642868e+02 2.36642868e+02 2.18642868e+02]
 [2.44571426e+02 2.37571426e+02 2.19571426e+02]]

...

[[1.41214325e+02 5.36428566e+01 5.27959023e+01]
 [1.37571487e+02 5.00000229e+01 4.82704086e+01]
 [1.39000046e+02 5.16428375e+01 4.90000000e+01]
 ...
 [2.50168350e+02 2.41168350e+02 2.24168350e+02]
 [2.50000000e+02 2.41000000e+02 2.24000000e+02]
 [2.49785736e+02 2.40785736e+02 2.23785736e+02]]

[[1.41974503e+02 5.39744949e+01 5.19744949e+01]
```

```
[1.40357147e+02 5.23571548e+01 5.03571548e+01]
[1.39642853e+02 5.16581650e+01 4.76122398e+01]
...
[2.51000000e+02 2.42000000e+02 2.25000000e+02]
[2.50933655e+02 2.41933655e+02 2.24933655e+02]
[2.50000000e+02 2.41000000e+02 2.24000000e+02]]

[[1.42571426e+02 5.45714302e+01 5.25714302e+01]
[1.42617371e+02 5.46173782e+01 5.26173782e+01]
[1.39642853e+02 5.18571434e+01 4.72142868e+01]
...
[2.51494919e+02 2.42494919e+02 2.25494919e+02]
[2.51000000e+02 2.42000000e+02 2.25000000e+02]
[2.51000000e+02 2.42000000e+02 2.25000000e+02]]]

[[[3.25867348e+01 3.25867348e+01 3.05867348e+01]
[3.20000000e+01 3.20000000e+01 3.20000000e+01]
[3.24285736e+01 3.20000000e+01 3.42142868e+01]
...
[3.66428566e+01 3.14285927e+01 2.90713844e+01]
[3.03570900e+01 2.53570900e+01 2.23570900e+01]
[2.30560646e+01 1.80560646e+01 1.50560656e+01]]]

[[3.14030609e+01 3.14030609e+01 2.94030609e+01]
[3.10714283e+01 3.10714283e+01 3.10714283e+01]
[3.24285736e+01 3.20000000e+01 3.42142868e+01]
...
[3.91428375e+01 3.39285736e+01 3.15713634e+01]
[3.35713768e+01 2.85713768e+01 2.55713768e+01]
[2.48570385e+01 1.98570385e+01 1.68570385e+01]]]

[[3.00000000e+01 3.00000000e+01 2.80000000e+01]
[3.07295914e+01 3.07295914e+01 3.07295914e+01]
[3.24285736e+01 3.20000000e+01 3.42142868e+01]
...
[4.04285278e+01 3.52142639e+01 3.28570557e+01]
[3.82142487e+01 3.32142487e+01 3.02142467e+01]
[2.99998608e+01 2.49998608e+01 2.19998608e+01]]]

...

[[2.54000000e+02 2.54000000e+02 2.54000000e+02]
[2.54000000e+02 2.54000000e+02 2.54000000e+02]
[2.54000000e+02 2.54000000e+02 2.54000000e+02]
...
[4.64285278e+01 2.27857361e+01 9.00000000e+00]
[4.59285583e+01 2.21428223e+01 9.71429443e+00]
[4.28570862e+01 1.90713501e+01 6.64282227e+00]]]

[[2.54000000e+02 2.54000000e+02 2.54000000e+02]
[2.54000000e+02 2.54000000e+02 2.54000000e+02]
[2.54000000e+02 2.54000000e+02 2.54000000e+02]
...
[4.63723946e+01 2.27296009e+01 8.94386578e+00]
[4.27856750e+01 1.87856750e+01 6.78567505e+00]
[3.95968971e+01 1.55968981e+01 3.59689808e+00]]]

[[2.54000000e+02 2.54000000e+02 2.54000000e+02]
[2.54000000e+02 2.54000000e+02 2.54000000e+02]
[2.54000000e+02 2.54000000e+02 2.54000000e+02]
...
[4.32090836e+01 1.87805557e+01 6.56629181e+00]
[4.10000000e+01 1.70000000e+01 5.00000000e+00]
[4.07143555e+01 1.67143555e+01 6.71435547e+00]]]

...

[[[2.35000000e+02 2.33000000e+02 1.84000000e+02]
[2.36928574e+02 2.34928574e+02 1.86928574e+02]
```



```
[2.38138702e+02 2.35924423e+02 1.88567276e+02]
...
[1.42204636e+02 1.06061691e+02 7.66331635e+01]
[1.36100983e+02 9.71009827e+01 6.61009827e+01]
[1.26105240e+02 8.51052399e+01 5.51052399e+01]]

[[2.35642853e+02 2.31642853e+02 1.83642853e+02]
[2.37874680e+02 2.33874680e+02 1.86874680e+02]
[2.38941971e+02 2.34941971e+02 1.88370544e+02]
...
[1.49318756e+02 1.13961540e+02 8.21509323e+01]
[1.34856155e+02 9.58561478e+01 6.48561478e+01]
[1.27820007e+02 8.68200073e+01 5.68200111e+01]]

[[2.35642853e+02 2.31642853e+02 1.83642853e+02]
[2.36945480e+02 2.32945480e+02 1.85182083e+02]
[2.37236603e+02 2.33236603e+02 1.86338013e+02]
...
[1.48766525e+02 1.13409317e+02 8.11950531e+01]
[1.38290176e+02 9.92901764e+01 6.67633896e+01]
[1.32347656e+02 9.13476562e+01 5.98208733e+01]]

...

[[2.55000000e+02 2.12133591e+02 1.86660324e+02]
[2.55000000e+02 2.14763367e+02 1.86763367e+02]
[2.53323334e+02 2.12650497e+02 1.84385498e+02]
...
[2.26525040e+02 1.69525040e+02 1.62525040e+02]
[2.52622208e+02 1.98511551e+02 1.89116196e+02]
[2.53545837e+02 2.02877548e+02 1.91211685e+02]]

[[2.52510468e+02 2.10364456e+02 1.82758881e+02]
[2.53116028e+02 2.12116028e+02 1.84116028e+02]
[2.53416138e+02 2.14300110e+02 1.85131409e+02]
...
[1.46211349e+02 8.82693710e+01 8.40953293e+01]
[2.23439270e+02 1.68443405e+02 1.61430969e+02]
[2.53081512e+02 2.00401398e+02 1.91860809e+02]]

[[2.52710815e+02 2.12458923e+02 1.86584869e+02]
[2.51163895e+02 2.12163895e+02 1.83163895e+02]
[2.47927551e+02 2.10776413e+02 1.81776413e+02]
...
[1.41158997e+02 8.71762772e+01 8.26849136e+01]
[1.58315674e+02 1.05315666e+02 9.93156662e+01]
[2.39812149e+02 1.89132050e+02 1.82106827e+02]]]

[[[2.23229599e+02 1.86229599e+02 8.22295914e+01]
[2.27816330e+02 1.90816330e+02 8.68163300e+01]
[2.31607132e+02 1.93607132e+02 9.26071396e+01]
...
[6.21326485e+01 2.38417950e+01 5.98976803e+00]
[6.32602386e+01 2.16173820e+01 4.85206604e+00]
[6.12141113e+01 1.95712547e+01 2.98453188e+00]]]

[[2.20500000e+02 1.83525497e+02 7.73775558e+01]
[2.20214294e+02 1.83214294e+02 7.89336777e+01]
[2.21770416e+02 1.83770416e+02 8.33112259e+01]
...
[6.63571014e+01 2.94132233e+01 3.19893765e+00]
[6.80765457e+01 2.81377678e+01 2.78571439e+00]
[6.53365631e+01 2.34079933e+01 5.96906424e-01]]]

[[2.15132660e+02 1.78994904e+02 6.78367386e+01]
[2.13642868e+02 1.76642868e+02 7.05714264e+01]
[2.16091843e+02 1.77923462e+02 7.95969391e+01]
...
[7.24439240e+01 3.48265762e+01 4.59136739e-02]
[7.98724670e+01 4.06734505e+01 3.87243915e+00]
[7.55713272e+01 3.37856102e+01 0.00000000e+00]]]
```

...

[ [2.47362259e+02 2.40362259e+02 1.86010223e+02]  
[2.47127548e+02 2.40127548e+02 1.85428574e+02]  
[2.45025513e+02 2.38214279e+02 1.81025528e+02]  
...  
[2.49000000e+02 2.41168381e+02 1.91663239e+02]  
[2.48729568e+02 2.40943863e+02 1.91300980e+02]  
[2.48000000e+02 2.40494934e+02 1.89857056e+02] ]

[ [2.50163315e+02 2.43163315e+02 1.91071503e+02]  
[2.48000015e+02 2.41000015e+02 1.88867401e+02]  
[2.47301056e+02 2.39673508e+02 1.85801071e+02]  
...  
[2.48071411e+02 2.41056107e+02 1.89102020e+02]  
[2.48000000e+02 2.41000000e+02 1.88867325e+02]  
[2.48331665e+02 2.41331665e+02 1.87423477e+02] ]

[ [2.50301025e+02 2.42301025e+02 1.93301025e+02]  
[2.46382660e+02 2.38714294e+02 1.88719376e+02]  
[2.47005112e+02 2.39362244e+02 1.86586731e+02]  
...  
[2.47214264e+02 2.40214264e+02 1.87204041e+02]  
[2.47071442e+02 2.40071442e+02 1.86071442e+02]  
[2.49000000e+02 2.42000000e+02 1.87229584e+02] ] ]

[ [ [6.97142868e+01 3.65000000e+01 3.78571415e+00]  
[7.13112259e+01 3.87397957e+01 4.09693861e+00]  
[6.95765305e+01 3.76479568e+01 2.29081607e+00]  
...  
[8.57296371e+01 7.10051575e+01 3.42143517e+01]  
[8.00407181e+01 6.71580582e+01 3.58723984e+01]  
[6.88163376e+01 5.69591942e+01 3.08164062e+01] ] ]

[ [7.57602081e+01 3.49030609e+01 2.83163333e+00]  
[7.48520432e+01 3.60663261e+01 2.99489808e+00]  
[6.94285660e+01 3.38571396e+01 5.61222248e-02]  
...  
[8.16378098e+01 5.68265762e+01 2.58265553e+01]  
[8.39948807e+01 5.93520241e+01 1.94999714e+01]  
[8.73317719e+01 6.36174850e+01 1.82603378e+01] ]

[ [6.67193832e+01 3.52193871e+01 2.21938777e+00]  
[6.24285660e+01 3.07142849e+01 1.53060742e-02]  
[5.93469391e+01 2.78316326e+01 9.18368548e-02]  
...  
[8.25203552e+01 6.37601357e+01 3.13570137e+01]  
[7.80000000e+01 5.59285469e+01 1.55867128e+01]  
[8.65716400e+01 6.06430664e+01 1.65002098e+01] ]

...

[ [2.00591873e+02 2.15158142e+02 2.14382751e+02]  
[1.96270523e+02 2.09841873e+02 2.02326660e+02]  
[1.92760223e+02 2.06545929e+02 1.90780655e+02]  
...  
[1.14596886e+02 1.18596886e+02 9.48316422e+01]  
[1.22800980e+02 1.27800980e+02 1.05800980e+02]  
[1.15928619e+02 1.20928619e+02 1.00928619e+02] ]

[ [1.94892899e+02 2.02750015e+02 2.14678635e+02]  
[2.08076462e+02 2.16928482e+02 2.22428528e+02]  
[1.96413361e+02 2.09556198e+02 2.00153214e+02]  
...  
[1.15285706e+02 1.18270401e+02 9.94592056e+01]  
[1.21846909e+02 1.26704025e+02 1.06775467e+02]  
[1.10285736e+02 1.15142853e+02 9.32142944e+01] ]

[ [1.85152374e+02 1.88152374e+02 2.06325897e+02]  
[2.01362289e+02 2.07959198e+02 2.18908234e+02]

```
[2.06392944e+02 2.17387833e+02 2.17765442e+02]
...
[1.17357178e+02 1.19709175e+02 1.05938766e+02]
[1.23499908e+02 1.26499908e+02 1.07499908e+02]
[1.07285645e+02 1.10285645e+02 8.92856445e+01]]]], shape=(32, 224, 224, 3), dtype=floa
t32) tf.Tensor(
[[0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0.]
 [0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0.]
 [0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]
 [0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]
 [0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]
 [0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0.]], shape=(32, 10), dtype=float32)
```

Notice how the image arrays come out as tensors of pixel values where as the labels come out as one-hot encodings (e.g. `[0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0.]` for `hamburger`).

## Model 0: Building a transfer learning model using the Keras Functional API

Alright, our data is tensor-ified, let's build a model.

To do so we're going to be using the `tf.keras.applications` module as it contains a series of already trained (on ImageNet) computer vision models as well as the Keras Functional API to construct our model.

We're going to go through the following steps:

1. Instantiate a pre-trained base model object by choosing a target model such as `EfficientNetB0` from `tf.keras.applications`, setting the `include_top` parameter to `False` (we do this because we're going to create our own top, which are the output layers for the model).
2. Set the base model's `trainable` attribute to `False` to freeze all of the weights in the pre-trained model.
3. Define an input layer for our model, for example, what shape of data should our model expect?
4. [Optional] Normalize the inputs to our model if it requires. Some computer vision models such as `ResNetV250` require their inputs to be between 0 & 1.

**Note:** As of writing, the `EfficientNet` models in the `tf.keras.applications` module do not require images to be normalized (pixel values between 0 and 1) on input, where as many of the other models do. I posted [an issue to the TensorFlow GitHub](#) about this and they confirmed this.

1. Pass the inputs to the base model.
2. Pool the outputs of the base model into a shape compatible with the output activation layer (turn base model output tensors into same shape as label tensor). This can be done using

output tensors into same shape as label tensors). This can be done using

`tf.keras.layers.GlobalAveragePooling2D()` or `tf.keras.layers.GlobalMaxPooling2D()`

though the former is more common in practice.

3. Create an output activation layer using `tf.keras.layers.Dense()` with the appropriate activation function and number of neurons.
4. Combine the inputs and outputs layer into a model using `tf.keras.Model()`.
5. Compile the model using the appropriate loss function and choose of optimizer.
6. Fit the model for desired number of epochs and with necessary callbacks (in our case, we'll start off with the TensorBoard callback).

Woah... that sounds like a lot. Before we get ahead of ourselves, let's see it in practice.

In [ ]:

```
# 1. Create base model with tf.keras.applications
base_model = tf.keras.applications.EfficientNetB0(include_top=False)

# 2. Freeze the base model (so the pre-learned patterns remain)
base_model.trainable = False

# 3. Create inputs into the base model
inputs = tf.keras.layers.Input(shape=(224, 224, 3), name="input_layer")

# 4. If using ResNet50V2, add this to speed up convergence, remove for EfficientNet
# x = tf.keras.layers.experimental.preprocessing.Rescaling(1./255)(inputs)

# 5. Pass the inputs to the base_model (note: using tf.keras.applications, EfficientNet i
nputs don't have to be normalized)
x = base_model(inputs)
# Check data shape after passing it to base_model
print(f"Shape after base_model: {x.shape}")

# 6. Average pool the outputs of the base model (aggregate all the most important informa
tion, reduce number of computations)
x = tf.keras.layers.GlobalAveragePooling2D(name="global_average_pooling_layer")(x)
print(f"After GlobalAveragePooling2D(): {x.shape}")

# 7. Create the output activation layer
outputs = tf.keras.layers.Dense(10, activation="softmax", name="output_layer")(x)

# 8. Combine the inputs with the outputs into a model
model_0 = tf.keras.Model(inputs, outputs)

# 9. Compile the model
model_0.compile(loss='categorical_crossentropy',
                optimizer=tf.keras.optimizers.Adam(),
                metrics=["accuracy"])

# 10. Fit the model (we use less steps for validation so it's faster)
history_10_percent = model_0.fit(train_data_10_percent,
                                epochs=5,
                                steps_per_epoch=len(train_data_10_percent),
                                validation_data=test_data_10_percent,
                                # Go through less of the validation data so epochs are
faster (we want faster experiments!)
                                validation_steps=int(0.25 * len(test_data_10_percent)),
                                # Track our model's training logs for visualization lat
er
                                callbacks=[create_tensorboard_callback("transfer_learning", "10_percent_feature_extract")])
```

Shape after base\_model: (None, 7, 7, 1280)

After GlobalAveragePooling2D(): (None, 1280)

Saving TensorBoard log files to: transfer\_learning/10\_percent\_feature\_extract/20210216-021515

Epoch 1/5

24/24 [=====] - 14s 313ms/step - loss: 2.1271 - accuracy: 0.2823  
- val\_loss: 1.3399 - val\_accuracy: 0.6743

Epoch 2/5

24/24 [=====] - 14s 313ms/step - loss: 1.8676 - accuracy: 0.6743

```
24/24 [=====] - 6s 223ms/step - loss: 1.2676 - accuracy: 0.6741
- val_loss: 0.8985 - val_accuracy: 0.8224
Epoch 3/5
24/24 [=====] - 6s 222ms/step - loss: 0.9113 - accuracy: 0.7824
- val_loss: 0.6983 - val_accuracy: 0.8520
Epoch 4/5
24/24 [=====] - 6s 224ms/step - loss: 0.7330 - accuracy: 0.8208
- val_loss: 0.6213 - val_accuracy: 0.8635
Epoch 5/5
24/24 [=====] - 6s 224ms/step - loss: 0.6232 - accuracy: 0.8519
- val_loss: 0.5716 - val_accuracy: 0.8651
```

**Nice! After a minute or so of training our model performs incredibly well on both the training (87%+ accuracy) and test sets (~83% accuracy).**

**This is incredible. All thanks to the power of transfer learning.**

**It's important to note the kind of transfer learning we used here is called feature extraction transfer learning, similar to what we did with the TensorFlow Hub models.**

**In other words, we passed our custom data to an already pre-trained model ( `EfficientNetB0` ), asked it "what patterns do you see?" and then put our own output layer on top to make sure the outputs were tailored to our desired number of classes.**

**We also used the Keras Functional API to build our model rather than the Sequential API. For now, the benefits of this main not seem clear but when you start to build more sophisticated models, you'll probably want to use the Functional API. So it's important to have exposure to this way of building models.**

▮ **Resource:** To see the benefits and use cases of the Functional API versus the Sequential API, check out the [TensorFlow Functional API documentation](#).

**Let's inspect the layers in our model, we'll start with the base.**

In [ ]:

```
# Check layers in our base model
for layer_number, layer in enumerate(base_model.layers):
    print(layer_number, layer.name)
```

```
0 input_1
1 rescaling
2 normalization
3 stem_conv_pad
4 stem_conv
5 stem_bn
6 stem_activation
7 block1a_dwconv
8 block1a_bn
9 block1a_activation
10 block1a_se_squeeze
11 block1a_se_reshape
12 block1a_se_reduce
13 block1a_se_expand
14 block1a_se_excite
15 block1a_project_conv
16 block1a_project_bn
17 block2a_expand_conv
18 block2a_expand_bn
19 block2a_expand_activation
20 block2a_dwconv_pad
21 block2a_dwconv
22 block2a_bn
23 block2a_activation
24 block2a_se_squeeze
25 block2a_se_reshape
26 block2a_se_reduce
27 block2a_se_expand
28 block2a_se_excite
29 block2a_project_conv
```

30 block2a\_project\_bn  
31 block2b\_expand\_conv  
32 block2b\_expand\_bn  
33 block2b\_expand\_activation  
34 block2b\_dwconv  
35 block2b\_bn  
36 block2b\_activation  
37 block2b\_se\_squeeze  
38 block2b\_se\_reshape  
39 block2b\_se\_reduce  
40 block2b\_se\_expand  
41 block2b\_se\_excite  
42 block2b\_project\_conv  
43 block2b\_project\_bn  
44 block2b\_drop  
45 block2b\_add  
46 block3a\_expand\_conv  
47 block3a\_expand\_bn  
48 block3a\_expand\_activation  
49 block3a\_dwconv\_pad  
50 block3a\_dwconv  
51 block3a\_bn  
52 block3a\_activation  
53 block3a\_se\_squeeze  
54 block3a\_se\_reshape  
55 block3a\_se\_reduce  
56 block3a\_se\_expand  
57 block3a\_se\_excite  
58 block3a\_project\_conv  
59 block3a\_project\_bn  
60 block3b\_expand\_conv  
61 block3b\_expand\_bn  
62 block3b\_expand\_activation  
63 block3b\_dwconv  
64 block3b\_bn  
65 block3b\_activation  
66 block3b\_se\_squeeze  
67 block3b\_se\_reshape  
68 block3b\_se\_reduce  
69 block3b\_se\_expand  
70 block3b\_se\_excite  
71 block3b\_project\_conv  
72 block3b\_project\_bn  
73 block3b\_drop  
74 block3b\_add  
75 block4a\_expand\_conv  
76 block4a\_expand\_bn  
77 block4a\_expand\_activation  
78 block4a\_dwconv\_pad  
79 block4a\_dwconv  
80 block4a\_bn  
81 block4a\_activation  
82 block4a\_se\_squeeze  
83 block4a\_se\_reshape  
84 block4a\_se\_reduce  
85 block4a\_se\_expand  
86 block4a\_se\_excite  
87 block4a\_project\_conv  
88 block4a\_project\_bn  
89 block4b\_expand\_conv  
90 block4b\_expand\_bn  
91 block4b\_expand\_activation  
92 block4b\_dwconv  
93 block4b\_bn  
94 block4b\_activation  
95 block4b\_se\_squeeze  
96 block4b\_se\_reshape  
97 block4b\_se\_reduce  
98 block4b\_se\_expand  
99 block4b\_se\_excite  
100 block4b\_project\_conv  
101 block4b\_project\_bn

102 block4b\_drop  
103 block4b\_add  
104 block4c\_expand\_conv  
105 block4c\_expand\_bn  
106 block4c\_expand\_activation  
107 block4c\_dwconv  
108 block4c\_bn  
109 block4c\_activation  
110 block4c\_se\_squeeze  
111 block4c\_se\_reshape  
112 block4c\_se\_reduce  
113 block4c\_se\_expand  
114 block4c\_se\_excite  
115 block4c\_project\_conv  
116 block4c\_project\_bn  
117 block4c\_drop  
118 block4c\_add  
119 block5a\_expand\_conv  
120 block5a\_expand\_bn  
121 block5a\_expand\_activation  
122 block5a\_dwconv  
123 block5a\_bn  
124 block5a\_activation  
125 block5a\_se\_squeeze  
126 block5a\_se\_reshape  
127 block5a\_se\_reduce  
128 block5a\_se\_expand  
129 block5a\_se\_excite  
130 block5a\_project\_conv  
131 block5a\_project\_bn  
132 block5b\_expand\_conv  
133 block5b\_expand\_bn  
134 block5b\_expand\_activation  
135 block5b\_dwconv  
136 block5b\_bn  
137 block5b\_activation  
138 block5b\_se\_squeeze  
139 block5b\_se\_reshape  
140 block5b\_se\_reduce  
141 block5b\_se\_expand  
142 block5b\_se\_excite  
143 block5b\_project\_conv  
144 block5b\_project\_bn  
145 block5b\_drop  
146 block5b\_add  
147 block5c\_expand\_conv  
148 block5c\_expand\_bn  
149 block5c\_expand\_activation  
150 block5c\_dwconv  
151 block5c\_bn  
152 block5c\_activation  
153 block5c\_se\_squeeze  
154 block5c\_se\_reshape  
155 block5c\_se\_reduce  
156 block5c\_se\_expand  
157 block5c\_se\_excite  
158 block5c\_project\_conv  
159 block5c\_project\_bn  
160 block5c\_drop  
161 block5c\_add  
162 block6a\_expand\_conv  
163 block6a\_expand\_bn  
164 block6a\_expand\_activation  
165 block6a\_dwconv\_pad  
166 block6a\_dwconv  
167 block6a\_bn  
168 block6a\_activation  
169 block6a\_se\_squeeze  
170 block6a\_se\_reshape  
171 block6a\_se\_reduce  
172 block6a\_se\_expand  
173 block6a\_se\_excite

```
174 block6a_project_conv
175 block6a_project_bn
176 block6b_expand_conv
177 block6b_expand_bn
178 block6b_expand_activation
179 block6b_dwconv
180 block6b_bn
181 block6b_activation
182 block6b_se_squeeze
183 block6b_se_reshape
184 block6b_se_reduce
185 block6b_se_expand
186 block6b_se_excite
187 block6b_project_conv
188 block6b_project_bn
189 block6b_drop
190 block6b_add
191 block6c_expand_conv
192 block6c_expand_bn
193 block6c_expand_activation
194 block6c_dwconv
195 block6c_bn
196 block6c_activation
197 block6c_se_squeeze
198 block6c_se_reshape
199 block6c_se_reduce
200 block6c_se_expand
201 block6c_se_excite
202 block6c_project_conv
203 block6c_project_bn
204 block6c_drop
205 block6c_add
206 block6d_expand_conv
207 block6d_expand_bn
208 block6d_expand_activation
209 block6d_dwconv
210 block6d_bn
211 block6d_activation
212 block6d_se_squeeze
213 block6d_se_reshape
214 block6d_se_reduce
215 block6d_se_expand
216 block6d_se_excite
217 block6d_project_conv
218 block6d_project_bn
219 block6d_drop
220 block6d_add
221 block7a_expand_conv
222 block7a_expand_bn
223 block7a_expand_activation
224 block7a_dwconv
225 block7a_bn
226 block7a_activation
227 block7a_se_squeeze
228 block7a_se_reshape
229 block7a_se_reduce
230 block7a_se_expand
231 block7a_se_excite
232 block7a_project_conv
233 block7a_project_bn
234 top_conv
235 top_bn
236 top_activation
```

**Wow, that's a lot of layers... to handcode all of those would've taken a fairly long time to do, yet we can still take advantage of them thanks to the power of transfer learning.**

**How about a summary of the base model?**

In [ ]:



```
base_model.summary()
```

Model: "efficientnetb0"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_1 (InputLayer)	[(None, None, None, 0		
rescaling (Rescaling)	(None, None, None, 3 0		input_1[0][0]
normalization (Normalization)	(None, None, None, 3 7		rescaling[0][0]
stem_conv_pad (ZeroPadding2D)	(None, None, None, 3 0		normalization[0][0]
stem_conv (Conv2D)	(None, None, None, 3 864		stem_conv_pad[0][0]
stem_bn (BatchNormalization)	(None, None, None, 3 128		stem_conv[0][0]
stem_activation (Activation)	(None, None, None, 3 0		stem_bn[0][0]
block1a_dwconv (DepthwiseConv2D)	(None, None, None, 3 288		stem_activation[0][0]
block1a_bn (BatchNormalization)	(None, None, None, 3 128		block1a_dwconv[0][0]
block1a_activation (Activation)	(None, None, None, 3 0		block1a_bn[0][0]
block1a_se_squeeze (GlobalAvera	(None, 32)	0	block1a_activation[0][0]
block1a_se_reshape (Reshape)	(None, 1, 1, 32)	0	block1a_se_squeeze[0][0]
block1a_se_reduce (Conv2D)	(None, 1, 1, 8)	264	block1a_se_reshape[0][0]
block1a_se_expand (Conv2D)	(None, 1, 1, 32)	288	block1a_se_reduce[0][0]
block1a_se_excite (Multiply)	(None, None, None, 3 0		block1a_activation[0][0] block1a_se_expand[0][0]
block1a_project_conv (Conv2D)	(None, None, None, 1 512		block1a_se_excite[0][0]

block1a_project_bn	(BatchNormal	(None, None, None, 1	64	block1a_project_conv[0][0]
block2a_expand_conv	(Conv2D)	(None, None, None, 9	1536	block1a_project_bn[0][0]
block2a_expand_bn	(BatchNormali	(None, None, None, 9	384	block2a_expand_conv[0][0]
block2a_expand_activation	(Acti	(None, None, None, 9	0	block2a_expand_bn[0][0]
block2a_dwconv_pad	(ZeroPadding	(None, None, None, 9	0	block2a_expand_activatio
block2a_dwconv	(DepthwiseConv2D	(None, None, None, 9	864	n[0][0]
block2a_bn	(BatchNormalization)	(None, None, None, 9	384	block2a_dwconv_pad[0][0]
block2a_activation	(Activation)	(None, None, None, 9	0	block2a_dwconv[0][0]
block2a_se_squeeze	(GlobalAvera	(None, 96)	0	block2a_activation[0][0]
block2a_se_reshape	(Reshape)	(None, 1, 1, 96)	0	block2a_se_squeeze[0][0]
block2a_se_reduce	(Conv2D)	(None, 1, 1, 4)	388	block2a_se_reshape[0][0]
block2a_se_expand	(Conv2D)	(None, 1, 1, 96)	480	block2a_se_reduce[0][0]
block2a_se_excite	(Multiply)	(None, None, None, 9	0	block2a_activation[0][0]
				block2a_se_expand[0][0]
block2a_project_conv	(Conv2D)	(None, None, None, 2	2304	block2a_se_excite[0][0]
block2a_project_bn	(BatchNormal	(None, None, None, 2	96	block2a_project_conv[0][0]
block2b_expand_conv	(Conv2D)	(None, None, None, 1	3456	block2a_project_bn[0][0]
block2b_expand_bn	(BatchNormali	(None, None, None, 1	576	block2b_expand_conv[0][0]

block2b_expand_activation	(Activation)	(None, None, None, 1 0	block2b_expand_bn[0][0]
block2b_dwconv	(DepthwiseConv2D)	(None, None, None, 1 1296	block2b_expand_activation[0][0]
block2b_bn	(BatchNormalization)	(None, None, None, 1 576	block2b_dwconv[0][0]
block2b_activation	(Activation)	(None, None, None, 1 0	block2b_bn[0][0]
block2b_se_squeeze	(GlobalAveragePooling2D)	(None, 144) 0	block2b_activation[0][0]
block2b_se_reshape	(Reshape)	(None, 1, 1, 144) 0	block2b_se_squeeze[0][0]
block2b_se_reduce	(Conv2D)	(None, 1, 1, 6) 870	block2b_se_reshape[0][0]
block2b_se_expand	(Conv2D)	(None, 1, 1, 144) 1008	block2b_se_reduce[0][0]
block2b_se_excite	(Multiply)	(None, None, None, 1 0	block2b_activation[0][0]
			block2b_se_expand[0][0]
block2b_project_conv	(Conv2D)	(None, None, None, 2 3456	block2b_se_excite[0][0]
block2b_project_bn	(BatchNormalization)	(None, None, None, 2 96	block2b_project_conv[0][0]
block2b_drop	(Dropout)	(None, None, None, 2 0	block2b_project_bn[0][0]
block2b_add	(Add)	(None, None, None, 2 0	block2b_drop[0][0]
			block2a_project_bn[0][0]
block3a_expand_conv	(Conv2D)	(None, None, None, 1 3456	block2b_add[0][0]
block3a_expand_bn	(BatchNormalization)	(None, None, None, 1 576	block3a_expand_conv[0][0]
block3a_expand_activation	(Activation)	(None, None, None, 1 0	block3a_expand_bn[0][0]
block3a_dwconv_pad	(ZeroPadding2D)	(None, None, None, 1 0	block3a_expand_activation[0][0]

block3a_dwconv (DepthwiseConv2D)	(None, None, None, 1 3600)		block3a_dwconv_pad[0][0]
block3a_bn (BatchNormalization)	(None, None, None, 1 576)		block3a_dwconv[0][0]
block3a_activation (Activation)	(None, None, None, 1 0)		block3a_bn[0][0]
block3a_se_squeeze (GlobalAveragePooling2D)	(None, 144)	0	block3a_activation[0][0]
block3a_se_reshape (Reshape)	(None, 1, 1, 144)	0	block3a_se_squeeze[0][0]
block3a_se_reduce (Conv2D)	(None, 1, 1, 6)	870	block3a_se_reshape[0][0]
block3a_se_expand (Conv2D)	(None, 1, 1, 144)	1008	block3a_se_reduce[0][0]
block3a_se_excite (Multiply)	(None, None, None, 1 0]		block3a_activation[0][0] block3a_se_expand[0][0]
block3a_project_conv (Conv2D)	(None, None, None, 4 5760)		block3a_se_excite[0][0]
block3a_project_bn (BatchNormalization)	(None, None, None, 4 160 0]		block3a_project_conv[0][0]
block3b_expand_conv (Conv2D)	(None, None, None, 2 9600)		block3a_project_bn[0][0]
block3b_expand_bn (BatchNormalization)	(None, None, None, 2 960]		block3b_expand_conv[0][0]
block3b_expand_activation (Activation)	(None, None, None, 2 0)		block3b_expand_bn[0][0]
block3b_dwconv (DepthwiseConv2D)	(None, None, None, 2 6000 n[0][0]		block3b_expand_activation[0][0]
block3b_bn (BatchNormalization)	(None, None, None, 2 960)		block3b_dwconv[0][0]
block3b_activation (Activation)	(None, None, None, 2 0)		block3b_bn[0][0]
block3b_se_squeeze (GlobalAveragePooling2D)	(None, 240)	0	block3b_activation[0][0]
block3b_se_reshape (Reshape)	(None, 1, 1, 240)	0	block3b_se_squeeze[0][0]

block3b_se_reduce (Conv2D)	(None, 1, 1, 10)	2410	block3b_se_reshape[0][0]
block3b_se_expand (Conv2D)	(None, 1, 1, 240)	2640	block3b_se_reduce[0][0]
block3b_se_excite (Multiply)	(None, None, None, 2 0]		block3b_activation[0][0] block3b_se_expand[0][0]
block3b_project_conv (Conv2D)	(None, None, None, 4 9600		block3b_se_excite[0][0]
block3b_project_bn (BatchNormal	(None, None, None, 4 160		block3b_project_conv[0][0]
block3b_drop (Dropout)	(None, None, None, 4 0]		block3b_project_bn[0][0]
block3b_add (Add)	(None, None, None, 4 0		block3b_drop[0][0] block3a_project_bn[0][0]
block4a_expand_conv (Conv2D)	(None, None, None, 2 9600		block3b_add[0][0]
block4a_expand_bn (BatchNormali	(None, None, None, 2 960		block4a_expand_conv[0][0]
block4a_expand_activation (Acti	(None, None, None, 2 0		block4a_expand_bn[0][0]
block4a_dwconv_pad (ZeroPadding	(None, None, None, 2 0		block4a_expand_activatio
block4a_dwconv (DepthwiseConv2D	(None, None, None, 2 2160		block4a_dwconv_pad[0][0]
block4a_bn (BatchNormalization)	(None, None, None, 2 960		block4a_dwconv[0][0]
block4a_activation (Activation)	(None, None, None, 2 0		block4a_bn[0][0]
block4a_se_squeeze (GlobalAvera	(None, 240)	0	block4a_activation[0][0]
block4a_se_reshape (Reshape)	(None, 1, 1, 240)	0	block4a_se_squeeze[0][0]
block4a_se_reduce (Conv2D)	(None, 1, 1, 10)	2410	block4a_se_reshape[0][0]

block4a_se_expand (Conv2D)	(None, 1, 1, 240)	2640	block4a_se_reduce[0][0]
block4a_se_excite (Multiply)	(None, None, None, 2 0 ]		block4a_activation[0][0] block4a_se_expand[0][0]
block4a_project_conv (Conv2D)	(None, None, None, 8 19200		block4a_se_excite[0][0]
block4a_project_bn (BatchNormal	(None, None, None, 8 320 0]		block4a_project_conv[0][0]
block4b_expand_conv (Conv2D)	(None, None, None, 4 38400		block4a_project_bn[0][0]
block4b_expand_bn (BatchNormali	(None, None, None, 4 1920 ]		block4b_expand_conv[0][0]
block4b_expand_activation (Acti	(None, None, None, 4 0		block4b_expand_bn[0][0]
block4b_dwconv (DepthwiseConv2D	(None, None, None, 4 4320 n[0][0]		block4b_expand_activatio
block4b_bn (BatchNormalization)	(None, None, None, 4 1920		block4b_dwconv[0][0]
block4b_activation (Activation)	(None, None, None, 4 0		block4b_bn[0][0]
block4b_se_squeeze (GlobalAvera	(None, 480)	0	block4b_activation[0][0]
block4b_se_reshape (Reshape)	(None, 1, 1, 480)	0	block4b_se_squeeze[0][0]
block4b_se_reduce (Conv2D)	(None, 1, 1, 20)	9620	block4b_se_reshape[0][0]
block4b_se_expand (Conv2D)	(None, 1, 1, 480)	10080	block4b_se_reduce[0][0]
block4b_se_excite (Multiply)	(None, None, None, 4 0 ]		block4b_activation[0][0] block4b_se_expand[0][0]
block4b_project_conv (Conv2D)	(None, None, None, 8 38400		block4b_se_excite[0][0]
block4b_project_bn (BatchNormal	(None, None, None, 8 320		block4b_project_conv[0][0]

0]			
block4b_drop (Dropout)	(None, None, None, 8 0		block4b_project_bn[0][0]
block4b_add (Add)	(None, None, None, 8 0		block4b_drop[0][0]
			block4a_project_bn[0][0]
block4c_expand_conv (Conv2D)	(None, None, None, 4 38400		block4b_add[0][0]
block4c_expand_bn (BatchNormali	(None, None, None, 4 1920		block4c_expand_conv[0][0]
block4c_expand_activation (Acti	(None, None, None, 4 0		block4c_expand_bn[0][0]
block4c_dwconv (DepthwiseConv2D	(None, None, None, 4 4320		block4c_expand_activation[0][0]
block4c_bn (BatchNormalization)	(None, None, None, 4 1920		block4c_dwconv[0][0]
block4c_activation (Activation)	(None, None, None, 4 0		block4c_bn[0][0]
block4c_se_squeeze (GlobalAvera	(None, 480)	0	block4c_activation[0][0]
block4c_se_reshape (Reshape)	(None, 1, 1, 480)	0	block4c_se_squeeze[0][0]
block4c_se_reduce (Conv2D)	(None, 1, 1, 20)	9620	block4c_se_reshape[0][0]
block4c_se_expand (Conv2D)	(None, 1, 1, 480)	10080	block4c_se_reduce[0][0]
block4c_se_excite (Multiply)	(None, None, None, 4 0		block4c_activation[0][0]
			block4c_se_expand[0][0]
block4c_project_conv (Conv2D)	(None, None, None, 8 38400		block4c_se_excite[0][0]
block4c_project_bn (BatchNormal	(None, None, None, 8 320		block4c_project_conv[0][0]
block4c_drop (Dropout)	(None, None, None, 8 0		block4c_project_bn[0][0]
block4c_add (Add)	(None, None, None, 8 0		block4c_drop[0][0]

block4b\_add[0][0]

block5a\_expand\_conv (Conv2D) (None, None, None, 4 38400

block4c\_add[0][0]

block5a\_expand\_bn (BatchNormali (None, None, None, 4 1920  
]

block5a\_expand\_conv[0][0]

block5a\_expand\_activation (Acti (None, None, None, 4 0

block5a\_expand\_bn[0][0]

block5a\_dwconv (DepthwiseConv2D (None, None, None, 4 12000  
n[0][0]

block5a\_expand\_activatio

block5a\_bn (BatchNormalization) (None, None, None, 4 1920

block5a\_dwconv[0][0]

block5a\_activation (Activation) (None, None, None, 4 0

block5a\_bn[0][0]

block5a\_se\_squeeze (GlobalAvera (None, 480) 0

block5a\_activation[0][0]

block5a\_se\_reshape (Reshape) (None, 1, 1, 480) 0

block5a\_se\_squeeze[0][0]

block5a\_se\_reduce (Conv2D) (None, 1, 1, 20) 9620

block5a\_se\_reshape[0][0]

block5a\_se\_expand (Conv2D) (None, 1, 1, 480) 10080

block5a\_se\_reduce[0][0]

block5a\_se\_excite (Multiply) (None, None, None, 4 0  
]  
]

block5a\_activation[0][0]

block5a\_se\_expand[0][0]

block5a\_project\_conv (Conv2D) (None, None, None, 1 53760

block5a\_se\_excite[0][0]

block5a\_project\_bn (BatchNormal (None, None, None, 1 448  
0]

block5a\_project\_conv[0][

block5b\_expand\_conv (Conv2D) (None, None, None, 6 75264

block5a\_project\_bn[0][0]

block5b\_expand\_bn (BatchNormali (None, None, None, 6 2688  
]

block5b\_expand\_conv[0][0]

block5b\_expand\_activation (Acti (None, None, None, 6 0

block5b\_expand\_bn[0][0]

block5b\_dwconv (DepthwiseConv2D (None, None, None, 6 16800

block5b\_expand\_activatio



n[0][0]

block5b_bn (BatchNormalization)	(None, None, None, 6 2688	block5b_dwconv[0][0]
block5b_activation (Activation)	(None, None, None, 6 0	block5b_bn[0][0]
block5b_se_squeeze (GlobalAvera	(None, 672) 0	block5b_activation[0][0]
block5b_se_reshape (Reshape)	(None, 1, 1, 672) 0	block5b_se_squeeze[0][0]
block5b_se_reduce (Conv2D)	(None, 1, 1, 28) 18844	block5b_se_reshape[0][0]
block5b_se_expand (Conv2D)	(None, 1, 1, 672) 19488	block5b_se_reduce[0][0]
block5b_se_excite (Multiply)	(None, None, None, 6 0	block5b_activation[0][0]
		block5b_se_expand[0][0]
block5b_project_conv (Conv2D)	(None, None, None, 1 75264	block5b_se_excite[0][0]
block5b_project_bn (BatchNormal	(None, None, None, 1 448	block5b_project_conv[0][0]
block5b_drop (Dropout)	(None, None, None, 1 0	block5b_project_bn[0][0]
block5b_add (Add)	(None, None, None, 1 0	block5b_drop[0][0]
		block5a_project_bn[0][0]
block5c_expand_conv (Conv2D)	(None, None, None, 6 75264	block5b_add[0][0]
block5c_expand_bn (BatchNormali	(None, None, None, 6 2688	block5c_expand_conv[0][0]
block5c_expand_activation (Acti	(None, None, None, 6 0	block5c_expand_bn[0][0]
block5c_dwconv (DepthwiseConv2D	(None, None, None, 6 16800	block5c_expand_activatio
n[0][0]		
block5c_bn (BatchNormalization)	(None, None, None, 6 2688	block5c_dwconv[0][0]
block5c_activation (Activation)	(None, None, None, 6 0	block5c_bn[0][0]

block5c_se_squeeze	(GlobalAveragePooling2D)	(None, 672)	0	block5c_activation[0][0]
block5c_se_reshape	(Reshape)	(None, 1, 1, 672)	0	block5c_se_squeeze[0][0]
block5c_se_reduce	(Conv2D)	(None, 1, 1, 28)	18844	block5c_se_reshape[0][0]
block5c_se_expand	(Conv2D)	(None, 1, 1, 672)	19488	block5c_se_reduce[0][0]
block5c_se_excite	(Multiply)	(None, None, None, 6 0]		block5c_activation[0][0] block5c_se_expand[0][0]
block5c_project_conv	(Conv2D)	(None, None, None, 1 75264		block5c_se_excite[0][0]
block5c_project_bn	(BatchNormalization)	(None, None, None, 1 448 0]		block5c_project_conv[0][0]
block5c_drop	(Dropout)	(None, None, None, 1 0]		block5c_project_bn[0][0]
block5c_add	(Add)	(None, None, None, 1 0		block5c_drop[0][0] block5b_add[0][0]
block6a_expand_conv	(Conv2D)	(None, None, None, 6 75264		block5c_add[0][0]
block6a_expand_bn	(BatchNormalization)	(None, None, None, 6 2688]		block6a_expand_conv[0][0]
block6a_expand_activation	(Activation)	(None, None, None, 6 0		block6a_expand_bn[0][0]
block6a_dwconv_pad	(ZeroPadding2D)	(None, None, None, 6 0 n[0][0]		block6a_expand_activation[0][0]
block6a_dwconv	(DepthwiseConv2D)	(None, None, None, 6 16800		block6a_dwconv_pad[0][0]
block6a_bn	(BatchNormalization)	(None, None, None, 6 2688		block6a_dwconv[0][0]
block6a_activation	(Activation)	(None, None, None, 6 0		block6a_bn[0][0]
block6a_se_squeeze	(GlobalAveragePooling2D)	(None, 672)	0	block6a_activation[0][0]

block6a_se_reshape (Reshape)	(None, 1, 1, 672)	0	block6a_se_squeeze[0][0]
block6a_se_reduce (Conv2D)	(None, 1, 1, 28)	18844	block6a_se_reshape[0][0]
block6a_se_expand (Conv2D)	(None, 1, 1, 672)	19488	block6a_se_reduce[0][0]
block6a_se_excite (Multiply)	(None, None, None, 6 0 ]		block6a_activation[0][0] block6a_se_expand[0][0]
block6a_project_conv (Conv2D)	(None, None, None, 1 129024		block6a_se_excite[0][0]
block6a_project_bn (BatchNormal 0]	(None, None, None, 1 768		block6a_project_conv[0][0]
block6b_expand_conv (Conv2D)	(None, None, None, 1 221184		block6a_project_bn[0][0]
block6b_expand_bn (BatchNormali ]	(None, None, None, 1 4608		block6b_expand_conv[0][0]
block6b_expand_activation (Acti	(None, None, None, 1 0		block6b_expand_bn[0][0]
block6b_dwconv (DepthwiseConv2D n[0][0]	(None, None, None, 1 28800		block6b_expand_activatio n[0][0]
block6b_bn (BatchNormalization)	(None, None, None, 1 4608		block6b_dwconv[0][0]
block6b_activation (Activation)	(None, None, None, 1 0		block6b_bn[0][0]
block6b_se_squeeze (GlobalAvera	(None, 1152)	0	block6b_activation[0][0]
block6b_se_reshape (Reshape)	(None, 1, 1, 1152)	0	block6b_se_squeeze[0][0]
block6b_se_reduce (Conv2D)	(None, 1, 1, 48)	55344	block6b_se_reshape[0][0]
block6b_se_expand (Conv2D)	(None, 1, 1, 1152)	56448	block6b_se_reduce[0][0]
block6b_se_excite (Multiply)	(None, None, None, 1 0 ]		block6b_activation[0][0] block6b_se_expand[0][0]

]		
block6b_project_conv (Conv2D)	(None, None, None, 1 221184	block6b_se_excite[0][0]
block6b_project_bn (BatchNormal	(None, None, None, 1 768	block6b_project_conv[0][0]
block6b_drop (Dropout)	(None, None, None, 1 0	block6b_project_bn[0][0]
block6b_add (Add)	(None, None, None, 1 0	block6b_drop[0][0]
		block6a_project_bn[0][0]
block6c_expand_conv (Conv2D)	(None, None, None, 1 221184	block6b_add[0][0]
block6c_expand_bn (BatchNormali	(None, None, None, 1 4608	block6c_expand_conv[0][0]
block6c_expand_activation (Acti	(None, None, None, 1 0	block6c_expand_bn[0][0]
block6c_dwconv (DepthwiseConv2D	(None, None, None, 1 28800	block6c_expand_activatio
n[0][0]		n[0][0]
block6c_bn (BatchNormalization)	(None, None, None, 1 4608	block6c_dwconv[0][0]
block6c_activation (Activation)	(None, None, None, 1 0	block6c_bn[0][0]
block6c_se_squeeze (GlobalAvera	(None, 1152) 0	block6c_activation[0][0]
block6c_se_reshape (Reshape)	(None, 1, 1, 1152) 0	block6c_se_squeeze[0][0]
block6c_se_reduce (Conv2D)	(None, 1, 1, 48) 55344	block6c_se_reshape[0][0]
block6c_se_expand (Conv2D)	(None, 1, 1, 1152) 56448	block6c_se_reduce[0][0]
block6c_se_excite (Multiply)	(None, None, None, 1 0	block6c_activation[0][0]
		block6c_se_expand[0][0]
block6c_project_conv (Conv2D)	(None, None, None, 1 221184	block6c_se_excite[0][0]
block6c_project_bn (BatchNormal	(None, None, None, 1 768	block6c_project_conv[0][0]

0]

block6c_drop (Dropout)	(None, None, None, 1 0	block6c_project_bn[0][0]
block6c_add (Add)	(None, None, None, 1 0	block6c_drop[0][0] block6b_add[0][0]
block6d_expand_conv (Conv2D)	(None, None, None, 1 221184	block6c_add[0][0]
block6d_expand_bn (BatchNormali	(None, None, None, 1 4608	block6d_expand_conv[0][0]
block6d_expand_activation (Acti	(None, None, None, 1 0	block6d_expand_bn[0][0]
block6d_dwconv (DepthwiseConv2D	(None, None, None, 1 28800	block6d_expand_activatio
n[0][0]		
block6d_bn (BatchNormalization)	(None, None, None, 1 4608	block6d_dwconv[0][0]
block6d_activation (Activation)	(None, None, None, 1 0	block6d_bn[0][0]
block6d_se_squeeze (GlobalAvera	(None, 1152)	0
		block6d_activation[0][0]
block6d_se_reshape (Reshape)	(None, 1, 1, 1152)	0
		block6d_se_squeeze[0][0]
block6d_se_reduce (Conv2D)	(None, 1, 1, 48)	55344
		block6d_se_reshape[0][0]
block6d_se_expand (Conv2D)	(None, 1, 1, 1152)	56448
		block6d_se_reduce[0][0]
block6d_se_excite (Multiply)	(None, None, None, 1 0	block6d_activation[0][0]
]		block6d_se_expand[0][0]
block6d_project_conv (Conv2D)	(None, None, None, 1 221184	block6d_se_excite[0][0]
block6d_project_bn (BatchNormal	(None, None, None, 1 768	block6d_project_conv[0][0]
0]		
block6d_drop (Dropout)	(None, None, None, 1 0	block6d_project_bn[0][0]
block6d_add (Add)	(None, None, None, 1 0	block6d_drop[0][0]

block6c\_add[0][0]

---

block7a\_expand\_conv (Conv2D) (None, None, None, 1 221184

block6d\_add[0][0]

---

block7a\_expand\_bn (BatchNormali (None, None, None, 1 4608  
]

block7a\_expand\_conv[0][0]

---

block7a\_expand\_activation (Acti (None, None, None, 1 0

block7a\_expand\_bn[0][0]

---

block7a\_dwconv (DepthwiseConv2D (None, None, None, 1 10368  
n[0][0]

block7a\_expand\_activatio

---

block7a\_bn (BatchNormalization) (None, None, None, 1 4608

block7a\_dwconv[0][0]

---

block7a\_activation (Activation) (None, None, None, 1 0

block7a\_bn[0][0]

---

block7a\_se\_squeeze (GlobalAvera (None, 1152) 0

block7a\_activation[0][0]

---

block7a\_se\_reshape (Reshape) (None, 1, 1, 1152) 0

block7a\_se\_squeeze[0][0]

---

block7a\_se\_reduce (Conv2D) (None, 1, 1, 48) 55344

block7a\_se\_reshape[0][0]

---

block7a\_se\_expand (Conv2D) (None, 1, 1, 1152) 56448

block7a\_se\_reduce[0][0]

---

block7a\_se\_excite (Multiply) (None, None, None, 1 0  
]  
]

block7a\_activation[0][0]

block7a\_se\_expand[0][0]

---

block7a\_project\_conv (Conv2D) (None, None, None, 3 368640

block7a\_se\_excite[0][0]

---

block7a\_project\_bn (BatchNormal (None, None, None, 3 1280  
0]

block7a\_project\_conv[0][

---

top\_conv (Conv2D) (None, None, None, 1 409600  
]

block7a\_project\_bn[0][0]

---

top\_bn (BatchNormalization) (None, None, None, 1 5120

top\_conv[0][0]

---

top\_activation (Activation) (None, None, None, 1 0

top\_bn[0][0]

---

=====  
Total params: 4,049,571

Trainable params: 0  
Non-trainable params: 4,049,571

You can see how each of the different layers have a certain number of parameters each. Since we are using a pre-trained model, you can think of all of these parameters are patterns the base model has learned on another dataset. And because we set `base_model.trainable = False`, these patterns remain as they are during training (they're frozen and don't get updated).

Alright that was the base model, let's see the summary of our overall model.

In [ ]:

```
# Check summary of model constructed with Functional API
model_0.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	[(None, 224, 224, 3)]	0
efficientnetb0 (Functional)	(None, None, None, 1280)	4049571
global_average_pooling_layer (None, 1280)		0
output_layer (Dense)	(None, 10)	12810
Total params: 4,062,381		
Trainable params: 12,810		
Non-trainable params: 4,049,571		

Our overall model has five layers but really, one of those layers ( `efficientnetb0` ) has 236 layers.



You can see how the output shape started out as `(None, 224, 224, 3)` for the input layer (the shape of our images) but was transformed to be `(None, 10)` by the output layer (the shape of our labels), where `None` is the placeholder for the batch size.

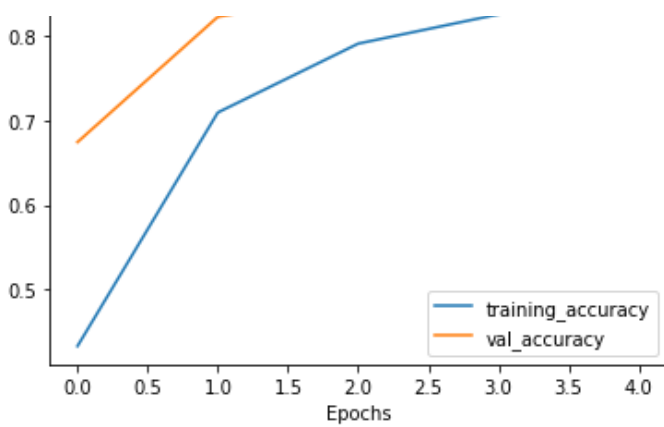
Notice too, the only trainable parameters in the model are those in the output layer.

How do our model's training curves look?

In [ ]:

```
# Check out our model's training curves
plot_loss_curves(history_10_percent)
```





## Getting a feature vector from a trained model

**Question:** What happens with the `tf.keras.layers.GlobalAveragePooling2D()` layer? I haven't seen it before.

The `tf.keras.layers.GlobalAveragePooling2D()` layer transforms a 4D tensor into a 2D tensor by averaging the values across the inner-axes.

The previous sentence is a bit of a mouthful, so let's see an example.

In [ ]:

```
# Define input tensor shape (same number of dimensions as the output of efficientnetb0)
input_shape = (1, 4, 4, 3)

# Create a random tensor
tf.random.set_seed(42)
input_tensor = tf.random.normal(input_shape)
print(f"Random input tensor:\n {input_tensor}\n")

# Pass the random tensor through a global average pooling 2D layer
global_average_pooled_tensor = tf.keras.layers.GlobalAveragePooling2D()(input_tensor)
print(f"2D global average pooled random tensor:\n {global_average_pooled_tensor}\n")

# Check the shapes of the different tensors
print(f"Shape of input tensor: {input_tensor.shape}")
print(f"Shape of 2D global averaged pooled input tensor: {global_average_pooled_tensor.shape}")
```

```
Random input tensor:
[[[ [ 0.3274685 -0.8426258  0.3194337 ]
    [-1.4075519 -2.3880599 -1.0392479 ]
    [-0.5573232  0.539707  1.6994323 ]
    [ 0.28893656 -1.5066116 -0.2645474 ]]]

 [[-0.59722406 -1.9171132 -0.62044144]
  [ 0.8504023 -0.40604794 -3.0258412 ]
  [ 0.9058464  0.29855987 -0.22561555]
  [-0.7616443 -1.891714 -0.9384712 ]]]

 [[ 0.77852213 -0.47338897  0.97772694]
  [ 0.24694404  0.20573747 -0.5256233 ]
  [ 0.32410017  0.02545409 -0.10638497]
  [-0.6369475  1.1603122  0.2507359 ]]]

 [[-0.41728497  0.40125778 -1.4145442 ]
  [-0.5931857 -1.6617213  0.33567193]
  [ 0.10815629  0.2347968 -0.56668764]
  [-0.35819843  0.88698614  0.52744764]]]]
```

```
2D global average pooled random tensor:
[[-0.09368646 -0.45840448 -0.2885598 ]]
```

```
Shape of input tensor: (1, 4, 4, 3)
```



Shape of 2D global averaged pooled input tensor: (1, 3)

You can see the `tf.keras.layers.GlobalAveragePooling2D()` layer condensed the input tensor from shape (1, 4, 4, 3) to (1, 3). It did so by averaging the `input_tensor` across the middle two axes.

We can replicate this operation using the `tf.reduce_mean()` operation and specifying the appropriate axes.

In [ ]:

```
# This is the same as GlobalAveragePooling2D()
tf.reduce_mean(input_tensor, axis=[1, 2]) # average across the middle axes
```

Out[ ]:

```
<tf.Tensor: shape=(1, 3), dtype=float32, numpy=array([[ -0.09368646, -0.45840448, -0.2885598 ]], dtype=float32)>
```

Doing this not only makes the output of the base model compatible with the input shape requirement of our output layer (`tf.keras.layers.Dense()`), it also condenses the information found by the base model into a lower dimension feature vector.

❏ **Note:** One of the reasons feature extraction transfer learning is named how it is is because what often happens is a pretrained model outputs a **feature vector** (a long tensor of numbers, in our case, this is the output of the `tf.keras.layers.GlobalAveragePooling2D()` layer) which can then be used to extract patterns out of.

❏ **Practice:** Do the same as the above cell but for `tf.keras.layers.GlobalMaxPool2D()`.

## Running a series of transfer learning experiments

We've seen the incredible results of transfer learning on 10% of the training data, what about 1% of the training data?

What kind of results do you think we can get using 100x less data than the original CNN models we built ourselves?

Why don't we answer that question while running the following modelling experiments:

1. `model_1`: Use feature extraction transfer learning on 1% of the training data with data augmentation.
2. `model_2`: Use feature extraction transfer learning on 10% of the training data with data augmentation.
3. `model_3`: Use fine-tuning transfer learning on 10% of the training data with data augmentation.
4. `model_4`: Use fine-tuning transfer learning on 100% of the training data with data augmentation.

While all of the experiments will be run on different versions of the training data, they will all be evaluated on the same test dataset, this ensures the results of each experiment are as comparable as possible.

All experiments will be done using the `EfficientNetB0` model within the `tf.keras.applications` module.

To make sure we're keeping track of our experiments, we'll use our `create_tensorboard_callback()` function to log all of the model training logs.

We'll construct each model using the Keras Functional API and instead of implementing data augmentation in the `ImageDataGenerator` class as we have previously, we're going to build it right into the model using the `tf.keras.layers.experimental.preprocessing` module.

Let's begin by downloading the data for experiment 1, using feature extraction transfer learning on 1% of the training data with data augmentation.

In [ ]:

```
# Download and unzip data
!wget https://storage.googleapis.com/ztm_tf_course/food_vision/10_food_classes_1_percent
```

```
.zip
unzip_data("10_food_classes_1_percent.zip")

# Create training and test dirs
train_dir_1_percent = "10_food_classes_1_percent/train/"
test_dir = "10_food_classes_1_percent/test/"
```

```
--2021-02-16 02:15:55-- https://storage.googleapis.com/ztm_tf_course/food_vision/10_food_classes_1_percent.zip
Resolving storage.googleapis.com (storage.googleapis.com)... 172.217.7.208, 172.217.9.208, 172.217.15.112, ...
Connecting to storage.googleapis.com (storage.googleapis.com)|172.217.7.208|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 133612354 (127M) [application/zip]
Saving to: '10_food_classes_1_percent.zip.1'

10_food_classes_1_p 100%[=====>] 127.42M 194MB/s in 0.7s

2021-02-16 02:15:56 (194 MB/s) - '10_food_classes_1_percent.zip.1' saved [133612354/133612354]
```

## How many images are we working with?

In [ ]:

```
# Walk through 1 percent data directory and list number of files
walk_through_dir("10_food_classes_1_percent")
```

```
There are 2 directories and 0 images in '10_food_classes_1_percent'.
There are 10 directories and 0 images in '10_food_classes_1_percent/train'.
There are 0 directories and 7 images in '10_food_classes_1_percent/train/ice_cream'.
There are 0 directories and 7 images in '10_food_classes_1_percent/train/ramen'.
There are 0 directories and 7 images in '10_food_classes_1_percent/train/chicken_wings'.
There are 0 directories and 7 images in '10_food_classes_1_percent/train/pizza'.
There are 0 directories and 7 images in '10_food_classes_1_percent/train/steak'.
There are 0 directories and 7 images in '10_food_classes_1_percent/train/fried_rice'.
There are 0 directories and 7 images in '10_food_classes_1_percent/train/hamburger'.
There are 0 directories and 7 images in '10_food_classes_1_percent/train/grilled_salmon'.
There are 0 directories and 7 images in '10_food_classes_1_percent/train/sushi'.
There are 0 directories and 7 images in '10_food_classes_1_percent/train/chicken_curry'.
There are 10 directories and 0 images in '10_food_classes_1_percent/test'.
There are 0 directories and 250 images in '10_food_classes_1_percent/test/ice_cream'.
There are 0 directories and 250 images in '10_food_classes_1_percent/test/ramen'.
There are 0 directories and 250 images in '10_food_classes_1_percent/test/chicken_wings'.
There are 0 directories and 250 images in '10_food_classes_1_percent/test/pizza'.
There are 0 directories and 250 images in '10_food_classes_1_percent/test/steak'.
There are 0 directories and 250 images in '10_food_classes_1_percent/test/fried_rice'.
There are 0 directories and 250 images in '10_food_classes_1_percent/test/hamburger'.
There are 0 directories and 250 images in '10_food_classes_1_percent/test/grilled_salmon'.
.
There are 0 directories and 250 images in '10_food_classes_1_percent/test/sushi'.
There are 0 directories and 250 images in '10_food_classes_1_percent/test/chicken_curry'.
```

Alright, looks like we've only got seven images of each class, this should be a bit of a challenge for our model.

**Note:** As with the 10% of data subset, the 1% of images were chosen at random from the original full training dataset. The test images are the same as the ones which have previously been used. If you want to see how this data was preprocessed, check out the [Food Vision Image Preprocessing notebook](#).

Time to load our images in as `tf.data.Dataset` objects, to do so, we'll use the `image_dataset_from_directory()` method.

In [ ]:

```
import tensorflow as tf
```

```
Found 70 files belonging to 10 classes.  
Found 2500 files belonging to 10 classes.
```

❗ **Note:** At the time of writing, the preprocessing layers we're using for data augmentation are in *experimental* status within the TensorFlow library. This means although the layers should be considered stable, the code may change slightly in a future version of TensorFlow. For more

information on the other preprocessing layers available and the different methods of data augmentation, check out the [Keras preprocessing layers guide](#) and the [TensorFlow data augmentation guide](#).

To use data augmentation right within our model we'll create a Keras Sequential model consisting of only data preprocessing layers, we can then use this Sequential model within another Functional model.

If that sounds confusing, it'll make sense once we create it in code.

The data augmentation transformations we're going to use are:

- [RandomFlip](#) - flips image on horizontal or vertical axis.
- [RandomRotation](#) - randomly rotates image by a specified amount.
- [RandomZoom](#) - randomly zooms into an image by specified amount.
- [RandomHeight](#) - randomly shifts image height by a specified amount.
- [RandomWidth](#) - randomly shifts image width by a specified amount.
- [Rescaling](#) - normalizes the image pixel values to be between 0 and 1, this is worth mentioning because it is required for some image models but since we're using the `tf.keras.applications` implementation of `EfficientNetB0`, it's not required.

There are more option but these will do for now.

In [ ]:

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.layers.experimental import preprocessing

# Create a data augmentation stage with horizontal flipping, rotations, zooms
data_augmentation = keras.Sequential([
    preprocessing.RandomFlip("horizontal"),
    preprocessing.RandomRotation(0.2),
    preprocessing.RandomZoom(0.2),
    preprocessing.RandomHeight(0.2),
    preprocessing.RandomWidth(0.2),
    # preprocessing.Rescaling(1./255) # keep for ResNet50V2, remove for EfficientNetB0
], name="data_augmentation")
```

And that's it! Our data augmentation Sequential model is ready to go. As you'll see shortly, we'll be able to slot this "model" as a layer into our transfer learning model later on.

But before we do that, let's test it out by passing random images through it.

In [ ]:

```
# View a random image
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import os
import random

target_class = random.choice(train_data_1_percent.class_names) # choose a random class
target_dir = "10_food_classes_1_percent/train/" + target_class # create the target directory
random_image = random.choice(os.listdir(target_dir)) # choose a random image from target directory
random_image_path = target_dir + "/" + random_image # create the choosen random image path
img = mpimg.imread(random_image_path) # read in the chosen target image
plt.imshow(img) # plot the target image
plt.title(f"Original random image from class: {target_class}")
plt.axis(False); # turn off the axes

# Augment the image
augmented_img = data_augmentation(tf.expand_dims(img, axis=0)) # data augmentation model requires shape (None, height, width, 3)
plt.figure()
```

```
plt.imshow(tf.squeeze(augmented_img)/255.) # requires normalization after augmentation
plt.title(f"Augmented random image from class: {target_class}")
plt.axis(False);
```

Original random image from class: grilled\_salmon



Augmented random image from class: grilled\_salmon



Run the cell above a few times and you can see the different random augmentations on different classes of images. Because we're going to add the data augmentation model as a layer in our upcoming transfer learning model, it'll apply these kind of random augmentations to each of the training images which passes through it.

Doing this will make our training dataset a little more varied. You can think of it as if you were taking a photo of food in real-life, not all of the images are going to be perfect, some of them are going to be orientated in strange ways. These are the kind of images we want our model to be able to handle.

Speaking of model, let's build one with the Functional API. We'll run through all of the same steps as before except for one difference, we'll add our data augmentation Sequential model as a layer immediately after the input layer.

## Model 1: Feature extraction transfer learning on 1% of the data with data augmentation

In [ ]:

```
# Setup input shape and base model, freezing the base model layers
input_shape = (224, 224, 3)
base_model = tf.keras.applications.EfficientNetB0(include_top=False)
base_model.trainable = False

# Create input layer
inputs = layers.Input(shape=input_shape, name="input_layer")

# Add in data augmentation Sequential model as a layer
x = data_augmentation(inputs)

# Give base_model inputs (after augmentation) and don't train it
x = base_model(x, training=False)

# Pool output features of base model
x = layers.GlobalAveragePooling2D(name="global_average_pooling_layer")(x)
```

```

# Put a dense layer on as the output
outputs = layers.Dense(10, activation="softmax", name="output_layer")(x)

# Make a model with inputs and outputs
model_1 = keras.Model(inputs, outputs)

# Compile the model
model_1.compile(loss="categorical_crossentropy",
                optimizer=tf.keras.optimizers.Adam(),
                metrics=["accuracy"])

# Fit the model
history_1_percent = model_1.fit(train_data_1_percent,
                                epochs=5,
                                steps_per_epoch=len(train_data_1_percent),
                                validation_data=test_data,
                                validation_steps=int(0.25* len(test_data)), # validate for less steps

                                # Track model training logs
                                callbacks=[create_tensorboard_callback("transfer_learning", "1_percent_data_aug")])

```

Saving TensorBoard log files to: transfer\_learning/1\_percent\_data\_aug/20210216-021736

Epoch 1/5  
3/3 [=====] - 10s 2s/step - loss: 2.4533 - accuracy: 0.0299 - val\_1\_loss: 2.2204 - val\_accuracy: 0.1645

Epoch 2/5  
3/3 [=====] - 4s 2s/step - loss: 2.1931 - accuracy: 0.0897 - val\_loss: 2.0827 - val\_accuracy: 0.2796

Epoch 3/5  
3/3 [=====] - 4s 2s/step - loss: 1.9338 - accuracy: 0.4006 - val\_loss: 1.9872 - val\_accuracy: 0.3536

Epoch 4/5  
3/3 [=====] - 4s 2s/step - loss: 1.8485 - accuracy: 0.4954 - val\_loss: 1.8808 - val\_accuracy: 0.4095

Epoch 5/5  
3/3 [=====] - 4s 2s/step - loss: 1.7186 - accuracy: 0.5845 - val\_loss: 1.7857 - val\_accuracy: 0.5099

**Wow! How cool is that? Using only 7 training images per class, using transfer learning our model was able to get ~40% accuracy on the validation set. This result is pretty amazing since the [original Food-101 paper](#) achieved 50.67% accuracy with all the data, namely, 750 training images per class (note: this metric was across 101 classes, not 10, we'll get to 101 classes soon).**

**If we check out a summary of our model, we should see the data augmentation layer just after the input layer.**

In [ ]:

```

# Check out model summary
model_1.summary()

```

Model: "model\_1"

Layer (type)	Output Shape	Param #
=====		
input_layer (InputLayer)	[(None, 224, 224, 3)]	0
-----		
data_augmentation (Sequentia	(None, None, None, 3)	0
-----		
efficientnetb0 (Functional)	(None, None, None, 1280)	4049571
-----		
global_average_pooling_layer	(None, 1280)	0
-----		
output_layer (Dense)	(None, 10)	12810
=====		
Total params: 4,062,381		
Trainable params: 12,810		
Non-trainable params: 4,049,571		



There it is. We've now got data augmentation built right into the our model. This means if we saved it and reloaded it somewhere else, the data augmentation layers would come with it.

The important thing to remember is **data augmentation only runs during training** . So if we were to evaluate or use our model for inference (predicting the class of an image) the data augmentation layers will be automatically turned off.

To see this in action, let's evaluate our model on the test data.

In [ ]:

```
# Evaluate on the test data
results_1_percent_data_aug = model_1.evaluate(test_data)
results_1_percent_data_aug
```

79/79 [=====] - 10s 122ms/step - loss: 1.8055 - accuracy: 0.4728

Out[ ]:


```
[1.8054912090301514, 0.47279998660087585]
```

The results here may be slightly better/worse than the log outputs of our model during training because during training we only evaluate our model on 25% of the test data using the line `validation_steps=int(0.25 * len(test_data))` . Doing this speeds up our epochs but still gives us enough of an idea of how our model is going.

Let's stay consistent and check out our model's loss curves.

In [ ]:

```
# How does the model go with a data augmentation layer with 1% of data
plot_loss_curves(history_1_percent)
```



It looks like the metrics on both datasets would improve if we kept training for more epochs. But we'll leave that

for now, we've got more experiments to do!

## Model 2: Feature extraction transfer learning with 10% of data and data augmentation

Alright, we've tested 1% of the training data with data augmentation, how about we try 10% of the data with data augmentation?

But wait...

□ **Question:** How do you know what experiments to run?

Great question.

The truth here is you often won't. Machine learning is still a very experimental practice. It's only after trying a fair few things that you'll start to develop an intuition of what to try.

My advice is to follow your curiosity as tenaciously as possible. If you feel like you want to try something, write the code for it and run it. See how it goes. The worst thing that'll happen is you'll figure out what doesn't work, the most valuable kind of knowledge.

From a practical standpoint, as we've talked about before, you'll want to reduce the amount of time between your initial experiments as much as possible. In other words, run a plethora of smaller experiments, using less data and less training iterations before you find something promising and then scale it up.

In the theme of scale, let's scale our 1% training data augmentation experiment up to 10% training data augmentation. That sentence doesn't really make sense but you get what I mean.

We're going to run through the exact same steps as the previous model, the only difference being using 10% of the training data instead of 1%.

In [ ]:

```
# Get 10% of the data of the 10 classes (uncomment if you haven't gotten "10_food_classes_10_percent.zip" already)
# !wget https://storage.googleapis.com/ztm_tf_course/food_vision/10_food_classes_10_percent.zip
# unzip_data("10_food_classes_10_percent.zip")

train_dir_10_percent = "10_food_classes_10_percent/train/"
test_dir = "10_food_classes_10_percent/test/"
```

Data downloaded. Let's create the dataloaders.

In [ ]:

```
# Setup data inputs
import tensorflow as tf
IMG_SIZE = (224, 224)
train_data_10_percent = tf.keras.preprocessing.image_dataset_from_directory(train_dir_10_percent,
                                                                              label_mode=
                                                                              "categorical",
                                                                              image_size=
                                                                              IMG_SIZE)
# Note: the test data is the same as the previous experiment, we could
# skip creating this, but we'll leave this here to practice.
test_data = tf.keras.preprocessing.image_dataset_from_directory(test_dir,
                                                                label_mode="categorical",
                                                                image_size=IMG_SIZE)
```

Found 750 files belonging to 10 classes.  
Found 2500 files belonging to 10 classes.

**Awesome! We've got 10x more images to work with, 75 per class instead of 7 per class.**



Let's build a model with data augmentation built in. We could reuse the data augmentation Sequential model we created before but we'll recreate it to practice.

In [ ]:

```
# Create a functional model with data augmentation
import tensorflow as tf
from tensorflow.keras import layers
from tensorflow.keras.layers.experimental import preprocessing
from tensorflow.keras.models import Sequential

# Build data augmentation layer
data_augmentation = Sequential([
    preprocessing.RandomFlip('horizontal'),
    preprocessing.RandomHeight(0.2),
    preprocessing.RandomWidth(0.2),
    preprocessing.RandomZoom(0.2),
    preprocessing.RandomRotation(0.2),
    # preprocessing.Rescaling(1./255) # keep for ResNet50V2, remove for EfficientNet
], name="data_augmentation")

# Setup the input shape to our model
input_shape = (224, 224, 3)

# Create a frozen base model
base_model = tf.keras.applications.EfficientNetB0(include_top=False)
base_model.trainable = False

# Create input and output layers
inputs = layers.Input(shape=input_shape, name="input_layer") # create input layer
x = data_augmentation(inputs) # augment our training images
x = base_model(x, training=False) # pass augmented images to base model but keep it in inference mode, so batchnorm layers don't get updated: https://keras.io/guides/transfer_learning/#build-a-model
x = layers.GlobalAveragePooling2D(name="global_average_pooling_layer")(x)
outputs = layers.Dense(10, activation="softmax", name="output_layer")(x)
model_2 = tf.keras.Model(inputs, outputs)

# Compile
model_2.compile(loss="categorical_crossentropy",
                optimizer=tf.keras.optimizers.Adam(lr=0.001), # use Adam optimizer with base learning rate
                metrics=["accuracy"])
```

## Creating a ModelCheckpoint callback

Our model is compiled and ready to be fit, so why haven't we fit it yet?

Well, for this experiment we're going to introduce a new callback, the `ModelCheckpoint` callback.

The `ModelCheckpoint` callback gives you the ability to save your model, as a whole in the `SavedModel` format or the [weights \(patterns\) only](#) to a specified directory as it trains.

This is helpful if you think your model is going to be training for a long time and you want to make backups of it as it trains. It also means if you think your model could benefit from being trained for longer, you can reload it from a specific checkpoint and continue training from there.

For example, say you fit a feature extraction transfer learning model for 5 epochs and you check the training curves and see it was still improving and you want to see if fine-tuning for another 5 epochs could help, you can load the checkpoint, unfreeze some (or all) of the base model layers and then continue training.

In fact, that's exactly what we're going to do.

But first, let's create a `ModelCheckpoint` callback. To do so, we have to specify a directory we'd like to save to.

In [ ]:

```
# Setup checkpoint path
checkpoint_path = "ten_percent_model_checkpoints_weights/checkpoint.ckpt" # note: remember saving directly to Colab is temporary

# Create a ModelCheckpoint callback that saves the model's weights only
checkpoint_callback = tf.keras.callbacks.ModelCheckpoint(filepath=checkpoint_path,
                                                         save_weights_only=True, # set to False to save the entire model
                                                         save_best_only=False, # set to True to save only the best model instead of a model every epoch
                                                         save_freq="epoch", # save every epoch
                                                         verbose=1)
```

**Question:** What's the difference between saving the entire model (SavedModel format) and saving the weights only?

The SavedModel format saves a model's architecture, weights and training configuration all in one folder. It makes it very easy to reload your model exactly how it is elsewhere. However, if you do not want to share all of these details with others, you may want to save and share the weights only (these will just be large tensors of non-human interpretable numbers). If disk space is an issue, saving the weights only is faster and takes up less space than saving the whole model.

Time to fit the model.

Because we're going to be fine-tuning it later, we'll create a variable `initial_epochs` and set it to 5 to use later.

We'll also add in our `checkpoint_callback` in our list of `callbacks`.

In [ ]:

```
# Fit the model saving checkpoints every epoch
initial_epochs = 5
history_10_percent_data_aug = model_2.fit(train_data_10_percent,
                                           epochs=initial_epochs,
                                           validation_data=test_data,
                                           validation_steps=int(0.25 * len(test_data)), # do less steps per validation (quicker)
                                           callbacks=[create_tensorboard_callback("transfer_learning", "10_percent_data_aug"),
                                           checkpoint_callback])
```

Saving TensorBoard log files to: transfer\_learning/10\_percent\_data\_aug/20210216-021854  
Epoch 1/5  
24/24 [=====] - 16s 452ms/step - loss: 2.1809 - accuracy: 0.2156  
- val\_loss: 1.5099 - val\_accuracy: 0.6299

Epoch 00001: saving model to ten\_percent\_model\_checkpoints\_weights/checkpoint.ckpt  
Epoch 2/5  
24/24 [=====] - 9s 358ms/step - loss: 1.4534 - accuracy: 0.6360  
- val\_loss: 1.0699 - val\_accuracy: 0.7582

Epoch 00002: saving model to ten\_percent\_model\_checkpoints\_weights/checkpoint.ckpt  
Epoch 3/5  
24/24 [=====] - 9s 375ms/step - loss: 1.0684 - accuracy: 0.7427  
- val\_loss: 0.8586 - val\_accuracy: 0.7878

Epoch 00003: saving model to ten\_percent\_model\_checkpoints\_weights/checkpoint.ckpt  
Epoch 4/5  
24/24 [=====] - 8s 338ms/step - loss: 0.9191 - accuracy: 0.7720  
- val\_loss: 0.7781 - val\_accuracy: 0.7993

Epoch 00004: saving model to ten\_percent\_model\_checkpoints\_weights/checkpoint.ckpt  
Epoch 5/5  
24/24 [=====] - 8s 332ms/step - loss: 0.7953 - accuracy: 0.7950

```
- val_loss: 0.7229 -val_accuracy: 0.7944
```

```
Epoch 00005: saving model to ten_percent_model_checkpoints_weights/checkpoint.ckpt
```

**Would you look at that! Looks like our `ModelCheckpoint` callback worked and our model saved its weights every epoch without too much overhead (saving the whole model takes longer than just the weights).**

**Let's evaluate our model and check its loss curves.**

```
In [ ]:
```

```
# Evaluate on the test data
results_10_percent_data_aug = model_2.evaluate(test_data)
results_10_percent_data_aug
```



```
79/79 [=====] - 10s 119ms/step - loss: 0.7047 - accuracy: 0.8080
```

```
Out[ ]:
```

```
[0.7046541571617126, 0.8080000281333923]
```

```
In [ ]:
```

```
# Plot model loss curves
plot_loss_curves(history_10_percent_data_aug)
```



Looking at these, our model's performance with 10% of the data and data augmentation isn't as good as the model with 10% of the data without data augmentation (see `model_0` results above), however the curves are trending in the right direction, meaning if we decided to train for longer, its metrics would likely improve.

Since we checkpointed (is that a word?) our model's weights, we might as well see what it's like to load it back in. We'll be able to test if it saved correctly by evaluating it on the test data.

To load saved model weights you can use the `load_weights()` method, passing it the path where your saved weights are stored.

```
# Load in saved model weights and evaluate model
model_2.load_weights(checkpoint_path)
loaded_weights_model_results = model_2.evaluate(test_data)
```

Now let's compare the results of our previously trained model and the loaded model. These results should very close if not exactly the same. The reason for minor differences comes down to the precision level of numbers calculated.

```
# If the results from our native model and the loaded weights are the same, this should output True
results_10_percent_data_aug == loaded_weights_model_results
```

If the above cell doesn't output `True`, it's because the numbers are close but not the *exact* same (due to how computers store numbers with degrees of precision).

```
import numpy as np
# Check to see if loaded model results are very close to native model results (should output True)
np.isclose(np.array(results_10_percent_data_aug), np.array(loaded_weights_model_results))
```

```
# Check the difference between the two results
print(np.array(results 10 percent data aug) - np.array(loaded weights model results))
```

Input data  
(10 classes of Food101)

Stays same  
(frozen, pre-trained on ImageNet)

Changes  
(unfrozen, gets fine-tuned on custom data)

Layers closer to the output layer get unfrozen/fine-tuned first

Bottom layers tend to stay frozen (or are last to get unfrozen)

Fully-connected (dense) classifier layer

**High-level example of fine-tuning an EfficientNet model. Bottom layers (layers closer to the input data) stay frozen whereas so top layers (layers closer to the output data) are updated during training.**

*frozen where as top layers (layers closer to the output data) are updated during training.*

So far our saved model has been trained using feature extraction transfer learning for 5 epochs on 10% of the training data and data augmentation.

This means all of the layers in the base model (EfficientNetB0) were frozen during training.

For our next experiment we're going to switch to fine-tuning transfer learning. This means we'll be using the same base model except we'll be unfreezing some of its layers (ones closest to the top) and running the model for a few more epochs.

The idea with fine-tuning is to start customizing the pre-trained model more to our own data.

❏ **Note:** Fine-tuning usually works best *after* training a feature extraction model for a few epochs and with large amounts of data. For more on this, check out [Keras' guide on Transfer learning & fine-tuning](#).

We've verified our loaded model's performance, let's check out its layers.

In [ ]:

```
# Layers in loaded model
model_2.layers
```

Out[ ]:

```
[<tensorflow.python.keras.engine.input_layer.InputLayer at 0x7fcec1c9208>,
<tensorflow.python.keras.engine.sequential.Sequential at 0x7fcec1c90b8>,
<tensorflow.python.keras.engine.functional.Functional at 0x7fcebe5bc630>,
<tensorflow.python.keras.layers.pooling.GlobalAveragePooling2D at 0x7fcebe62b8d0>,
<tensorflow.python.keras.layers.core.Dense at 0x7fcebe5f1e80>]
```

In [ ]:

```
for layer in model_2.layers:
    print(layer.trainable)
```

```
True
True
False
True
True
```

Looking good. We've got an input layer, a Sequential layer (the data augmentation model), a Functional layer (EfficientNetB0), a pooling layer and a Dense layer (the output layer).

How about a summary?

In [ ]:

```
model_2.summary()
```

Model: "model\_2"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	[(None, 224, 224, 3)]	0
data_augmentation (Sequentia	(None, None, None, 3)	0
efficientnetb0 (Functional)	(None, None, None, 1280)	4049571
global_average_pooling_layer	(None, 1280)	0
output_layer (Dense)	(None, 10)	12810
Total params: 4,062,381		
Trainable params: 12,810		
Non-trainable params: 4,049,571		

Alright, it looks like all of the layers in the `efficientnetb0` layer are frozen. We can confirm this using the `trainable_variables` attribute.

In [ ]:

```
# How many layers are trainable in our base model?
print(len(model_2.layers[2].trainable_variables)) # layer at index 2 is the EfficientNetB
0 layer (the base model)
```

0

This is the same as our base model.

In [ ]:

```
print(len(base_model.trainable_variables))
```

0

We can even check layer by layer to see if they're trainable.

In [ ]:

```
# Check which layers are tuneable (trainable)
for layer_number, layer in enumerate(base_model.layers):
    print(layer_number, layer.name, layer.trainable)
```

```
0 input_3 False
1 rescaling_2 False
2 normalization_2 False
3 stem_conv_pad False
4 stem_conv False
5 stem_bn False
6 stem_activation False
7 block1a_dwconv False
8 block1a_bn False
9 block1a_activation False
10 block1a_se_squeeze False
11 block1a_se_reshape False
12 block1a_se_reduce False
13 block1a_se_expand False
14 block1a_se_excite False
15 block1a_project_conv False
16 block1a_project_bn False
17 block2a_expand_conv False
18 block2a_expand_bn False
19 block2a_expand_activation False
20 block2a_dwconv_pad False
21 block2a_dwconv False
22 block2a_bn False
23 block2a_activation False
24 block2a_se_squeeze False
25 block2a_se_reshape False
26 block2a_se_reduce False
27 block2a_se_expand False
28 block2a_se_excite False
29 block2a_project_conv False
30 block2a_project_bn False
31 block2b_expand_conv False
32 block2b_expand_bn False
33 block2b_expand_activation False
34 block2b_dwconv False
35 block2b_bn False
36 block2b_activation False
37 block2b_se_squeeze False
38 block2b_se_reshape False
39 block2b_se_reduce False
40 block2b_se_expand False
41 block2b_se_excite False
```

41 block2b\_se\_excite False  
42 block2b\_project\_conv False  
43 block2b\_project\_bn False  
44 block2b\_drop False  
45 block2b\_add False  
46 block3a\_expand\_conv False  
47 block3a\_expand\_bn False  
48 block3a\_expand\_activation False  
49 block3a\_dwconv\_pad False  
50 block3a\_dwconv False  
51 block3a\_bn False  
52 block3a\_activation False  
53 block3a\_se\_squeeze False  
54 block3a\_se\_reshape False  
55 block3a\_se\_reduce False  
56 block3a\_se\_expand False  
57 block3a\_se\_excite False  
58 block3a\_project\_conv False  
59 block3a\_project\_bn False  
60 block3b\_expand\_conv False  
61 block3b\_expand\_bn False  
62 block3b\_expand\_activation False  
63 block3b\_dwconv False  
64 block3b\_bn False  
65 block3b\_activation False  
66 block3b\_se\_squeeze False  
67 block3b\_se\_reshape False  
68 block3b\_se\_reduce False  
69 block3b\_se\_expand False  
70 block3b\_se\_excite False  
71 block3b\_project\_conv False  
72 block3b\_project\_bn False  
73 block3b\_drop False  
74 block3b\_add False  
75 block4a\_expand\_conv False  
76 block4a\_expand\_bn False  
77 block4a\_expand\_activation False  
78 block4a\_dwconv\_pad False  
79 block4a\_dwconv False  
80 block4a\_bn False  
81 block4a\_activation False  
82 block4a\_se\_squeeze False  
83 block4a\_se\_reshape False  
84 block4a\_se\_reduce False  
85 block4a\_se\_expand False  
86 block4a\_se\_excite False  
87 block4a\_project\_conv False  
88 block4a\_project\_bn False  
89 block4b\_expand\_conv False  
90 block4b\_expand\_bn False  
91 block4b\_expand\_activation False  
92 block4b\_dwconv False  
93 block4b\_bn False  
94 block4b\_activation False  
95 block4b\_se\_squeeze False  
96 block4b\_se\_reshape False  
97 block4b\_se\_reduce False  
98 block4b\_se\_expand False  
99 block4b\_se\_excite False  
100 block4b\_project\_conv False  
101 block4b\_project\_bn False  
102 block4b\_drop False  
103 block4b\_add False  
104 block4c\_expand\_conv False  
105 block4c\_expand\_bn False  
106 block4c\_expand\_activation False  
107 block4c\_dwconv False  
108 block4c\_bn False  
109 block4c\_activation False  
110 block4c\_se\_squeeze False  
111 block4c\_se\_reshape False  
112 block4c\_se\_reduce False  
113 block4c\_se\_expand False  
114 block4c\_se\_excite False

113 block4c\_se\_expand False  
114 block4c\_se\_excite False  
115 block4c\_project\_conv False  
116 block4c\_project\_bn False  
117 block4c\_drop False  
118 block4c\_add False  
119 block5a\_expand\_conv False  
120 block5a\_expand\_bn False  
121 block5a\_expand\_activation False  
122 block5a\_dwconv False  
123 block5a\_bn False  
124 block5a\_activation False  
125 block5a\_se\_squeeze False  
126 block5a\_se\_reshape False  
127 block5a\_se\_reduce False  
128 block5a\_se\_expand False  
129 block5a\_se\_excite False  
130 block5a\_project\_conv False  
131 block5a\_project\_bn False  
132 block5b\_expand\_conv False  
133 block5b\_expand\_bn False  
134 block5b\_expand\_activation False  
135 block5b\_dwconv False  
136 block5b\_bn False  
137 block5b\_activation False  
138 block5b\_se\_squeeze False  
139 block5b\_se\_reshape False  
140 block5b\_se\_reduce False  
141 block5b\_se\_expand False  
142 block5b\_se\_excite False  
143 block5b\_project\_conv False  
144 block5b\_project\_bn False  
145 block5b\_drop False  
146 block5b\_add False  
147 block5c\_expand\_conv False  
148 block5c\_expand\_bn False  
149 block5c\_expand\_activation False  
150 block5c\_dwconv False  
151 block5c\_bn False  
152 block5c\_activation False  
153 block5c\_se\_squeeze False  
154 block5c\_se\_reshape False  
155 block5c\_se\_reduce False  
156 block5c\_se\_expand False  
157 block5c\_se\_excite False  
158 block5c\_project\_conv False  
159 block5c\_project\_bn False  
160 block5c\_drop False  
161 block5c\_add False  
162 block6a\_expand\_conv False  
163 block6a\_expand\_bn False  
164 block6a\_expand\_activation False  
165 block6a\_dwconv\_pad False  
166 block6a\_dwconv False  
167 block6a\_bn False  
168 block6a\_activation False  
169 block6a\_se\_squeeze False  
170 block6a\_se\_reshape False  
171 block6a\_se\_reduce False  
172 block6a\_se\_expand False  
173 block6a\_se\_excite False  
174 block6a\_project\_conv False  
175 block6a\_project\_bn False  
176 block6b\_expand\_conv False  
177 block6b\_expand\_bn False  
178 block6b\_expand\_activation False  
179 block6b\_dwconv False  
180 block6b\_bn False  
181 block6b\_activation False  
182 block6b\_se\_squeeze False  
183 block6b\_se\_reshape False  
184 block6b\_se\_reduce False  
185 block6b\_se\_excite False



```
185 block6b_se_expand False
186 block6b_se_excite False
187 block6b_project_conv False
188 block6b_project_bn False
189 block6b_drop False
190 block6b_add False
191 block6c_expand_conv False
192 block6c_expand_bn False
193 block6c_expand_activation False
194 block6c_dwconv False
195 block6c_bn False
196 block6c_activation False
197 block6c_se_squeeze False
198 block6c_se_reshape False
199 block6c_se_reduce False
200 block6c_se_expand False
201 block6c_se_excite False
202 block6c_project_conv False
203 block6c_project_bn False
204 block6c_drop False
205 block6c_add False
206 block6d_expand_conv False
207 block6d_expand_bn False
208 block6d_expand_activation False
209 block6d_dwconv False
210 block6d_bn False
211 block6d_activation False
212 block6d_se_squeeze False
213 block6d_se_reshape False
214 block6d_se_reduce False
215 block6d_se_expand False
216 block6d_se_excite False
217 block6d_project_conv False
218 block6d_project_bn False
219 block6d_drop False
220 block6d_add False
221 block7a_expand_conv False
222 block7a_expand_bn False
223 block7a_expand_activation False
224 block7a_dwconv False
225 block7a_bn False
226 block7a_activation False
227 block7a_se_squeeze False
228 block7a_se_reshape False
229 block7a_se_reduce False
230 block7a_se_expand False
231 block7a_se_excite False
232 block7a_project_conv False
233 block7a_project_bn False
234 top_conv False
235 top_bn False
236 top_activation False
```

**Beautiful. This is exactly what we're after.**

**Now to fine-tune the base model to our own data, we're going to unfreeze the top 10 layers and continue training our model for another 5 epochs.**

**This means all of the base model's layers except for the last 10 will remain frozen and untrainable. And the weights in the remaining unfrozen layers will be updated during training.**

**Ideally, we should see the model's performance improve.**

□ **Question:** How many layers should you unfreeze when training?

**There's no set rule for this. You could unfreeze every layer in the pretrained model or you could try unfreezing one layer at a time. Best to experiment with different amounts of unfreezing and fine-tuning to see what happens. Generally, the less data you have, the less layers you want to unfreeze and the more gradually you want to fine-tune.**

▮ **Resource:** The [ULMFiT \(Universal Language Model Fine-tuning for Text Classification\) paper](#) has a great series of experiments on fine-tuning models.

To begin fine-tuning, we'll unfreeze the entire base model by setting its `trainable` attribute to `True`. Then we'll refreeze every layer in the base model except for the last 10 by looping through them and setting their `trainable` attribute to `False`. Finally, we'll recompile the model.

In [ ]:

```
base_model.trainable = True

# Freeze all layers except for the
for layer in base_model.layers[:-10]:
    layer.trainable = False

# Recompile the model (always recompile after any adjustments to a model)
model_2.compile(loss="categorical_crossentropy",
                 optimizer=tf.keras.optimizers.Adam(lr=0.0001), # lr is 10x lower than before for fine-tuning
                 metrics=["accuracy"])
```

Wonderful, now let's check which layers of the pretrained model are trainable.

In [ ]:

```
# Check which layers are tuneable (trainable)
for layer_number, layer in enumerate(base_model.layers):
    print(layer_number, layer.name, layer.trainable)
```

```
0 input_3 False
1 rescaling_2 False
2 normalization_2 False
3 stem_conv_pad False
4 stem_conv False
5 stem_bn False
6 stem_activation False
7 block1a_dwconv False
8 block1a_bn False
9 block1a_activation False
10 block1a_se_squeeze False
11 block1a_se_reshape False
12 block1a_se_reduce False
13 block1a_se_expand False
14 block1a_se_excite False
15 block1a_project_conv False
16 block1a_project_bn False
17 block2a_expand_conv False
18 block2a_expand_bn False
19 block2a_expand_activation False
20 block2a_dwconv_pad False
21 block2a_dwconv False
22 block2a_bn False
23 block2a_activation False
24 block2a_se_squeeze False
25 block2a_se_reshape False
26 block2a_se_reduce False
27 block2a_se_expand False
28 block2a_se_excite False
29 block2a_project_conv False
30 block2a_project_bn False
31 block2b_expand_conv False
32 block2b_expand_bn False
33 block2b_expand_activation False
34 block2b_dwconv False
35 block2b_bn False
36 block2b_activation False
37 block2b_se_squeeze False
38 block2b_se_reshape False
39 block2b se reduce False
```

40 block2b\_se\_expand False  
41 block2b\_se\_excite False  
42 block2b\_project\_conv False  
43 block2b\_project\_bn False  
44 block2b\_drop False  
45 block2b\_add False  
46 block3a\_expand\_conv False  
47 block3a\_expand\_bn False  
48 block3a\_expand\_activation False  
49 block3a\_dwconv\_pad False  
50 block3a\_dwconv False  
51 block3a\_bn False  
52 block3a\_activation False  
53 block3a\_se\_squeeze False  
54 block3a\_se\_reshape False  
55 block3a\_se\_reduce False  
56 block3a\_se\_expand False  
57 block3a\_se\_excite False  
58 block3a\_project\_conv False  
59 block3a\_project\_bn False  
60 block3b\_expand\_conv False  
61 block3b\_expand\_bn False  
62 block3b\_expand\_activation False  
63 block3b\_dwconv False  
64 block3b\_bn False  
65 block3b\_activation False  
66 block3b\_se\_squeeze False  
67 block3b\_se\_reshape False  
68 block3b\_se\_reduce False  
69 block3b\_se\_expand False  
70 block3b\_se\_excite False  
71 block3b\_project\_conv False  
72 block3b\_project\_bn False  
73 block3b\_drop False  
74 block3b\_add False  
75 block4a\_expand\_conv False  
76 block4a\_expand\_bn False  
77 block4a\_expand\_activation False  
78 block4a\_dwconv\_pad False  
79 block4a\_dwconv False  
80 block4a\_bn False  
81 block4a\_activation False  
82 block4a\_se\_squeeze False  
83 block4a\_se\_reshape False  
84 block4a\_se\_reduce False  
85 block4a\_se\_expand False  
86 block4a\_se\_excite False  
87 block4a\_project\_conv False  
88 block4a\_project\_bn False  
89 block4b\_expand\_conv False  
90 block4b\_expand\_bn False  
91 block4b\_expand\_activation False  
92 block4b\_dwconv False  
93 block4b\_bn False  
94 block4b\_activation False  
95 block4b\_se\_squeeze False  
96 block4b\_se\_reshape False  
97 block4b\_se\_reduce False  
98 block4b\_se\_expand False  
99 block4b\_se\_excite False  
100 block4b\_project\_conv False  
101 block4b\_project\_bn False  
102 block4b\_drop False  
103 block4b\_add False  
104 block4c\_expand\_conv False  
105 block4c\_expand\_bn False  
106 block4c\_expand\_activation False  
107 block4c\_dwconv False  
108 block4c\_bn False  
109 block4c\_activation False  
110 block4c\_se\_squeeze False  
111 block4c\_se\_reshape False

112 block4c\_se\_reduce False  
113 block4c\_se\_expand False  
114 block4c\_se\_excite False  
115 block4c\_project\_conv False  
116 block4c\_project\_bn False  
117 block4c\_drop False  
118 block4c\_add False  
119 block5a\_expand\_conv False  
120 block5a\_expand\_bn False  
121 block5a\_expand\_activation False  
122 block5a\_dwconv False  
123 block5a\_bn False  
124 block5a\_activation False  
125 block5a\_se\_squeeze False  
126 block5a\_se\_reshape False  
127 block5a\_se\_reduce False  
128 block5a\_se\_expand False  
129 block5a\_se\_excite False  
130 block5a\_project\_conv False  
131 block5a\_project\_bn False  
132 block5b\_expand\_conv False  
133 block5b\_expand\_bn False  
134 block5b\_expand\_activation False  
135 block5b\_dwconv False  
136 block5b\_bn False  
137 block5b\_activation False  
138 block5b\_se\_squeeze False  
139 block5b\_se\_reshape False  
140 block5b\_se\_reduce False  
141 block5b\_se\_expand False  
142 block5b\_se\_excite False  
143 block5b\_project\_conv False  
144 block5b\_project\_bn False  
145 block5b\_drop False  
146 block5b\_add False  
147 block5c\_expand\_conv False  
148 block5c\_expand\_bn False  
149 block5c\_expand\_activation False  
150 block5c\_dwconv False  
151 block5c\_bn False  
152 block5c\_activation False  
153 block5c\_se\_squeeze False  
154 block5c\_se\_reshape False  
155 block5c\_se\_reduce False  
156 block5c\_se\_expand False  
157 block5c\_se\_excite False  
158 block5c\_project\_conv False  
159 block5c\_project\_bn False  
160 block5c\_drop False  
161 block5c\_add False  
162 block6a\_expand\_conv False  
163 block6a\_expand\_bn False  
164 block6a\_expand\_activation False  
165 block6a\_dwconv\_pad False  
166 block6a\_dwconv False  
167 block6a\_bn False  
168 block6a\_activation False  
169 block6a\_se\_squeeze False  
170 block6a\_se\_reshape False  
171 block6a\_se\_reduce False  
172 block6a\_se\_expand False  
173 block6a\_se\_excite False  
174 block6a\_project\_conv False  
175 block6a\_project\_bn False  
176 block6b\_expand\_conv False  
177 block6b\_expand\_bn False  
178 block6b\_expand\_activation False  
179 block6b\_dwconv False  
180 block6b\_bn False  
181 block6b\_activation False  
182 block6b\_se\_squeeze False  
183 block6b\_se\_reshape False

```
184 block6b_se_reduce False
185 block6b_se_expand False
186 block6b_se_excite False
187 block6b_project_conv False
188 block6b_project_bn False
189 block6b_drop False
190 block6b_add False
191 block6c_expand_conv False
192 block6c_expand_bn False
193 block6c_expand_activation False
194 block6c_dwconv False
195 block6c_bn False
196 block6c_activation False
197 block6c_se_squeeze False
198 block6c_se_reshape False
199 block6c_se_reduce False
200 block6c_se_expand False
201 block6c_se_excite False
202 block6c_project_conv False
203 block6c_project_bn False
204 block6c_drop False
205 block6c_add False
206 block6d_expand_conv False
207 block6d_expand_bn False
208 block6d_expand_activation False
209 block6d_dwconv False
210 block6d_bn False
211 block6d_activation False
212 block6d_se_squeeze False
213 block6d_se_reshape False
214 block6d_se_reduce False
215 block6d_se_expand False
216 block6d_se_excite False
217 block6d_project_conv False
218 block6d_project_bn False
219 block6d_drop False
220 block6d_add False
221 block7a_expand_conv False
222 block7a_expand_bn False
223 block7a_expand_activation False
224 block7a_dwconv False
225 block7a_bn False
226 block7a_activation False
227 block7a_se_squeeze True
228 block7a_se_reshape True
229 block7a_se_reduce True
230 block7a_se_expand True
231 block7a_se_excite True
232 block7a_project_conv True
233 block7a_project_bn True
234 top_conv True
235 top_bn True
236 top_activation True
```

**Nice! It seems all layers except for the last 10 are frozen and untrainable. This means only the last 10 layers of the base model along with the output layer will have their weights updated during training.**

□ **Question:** Why did we recompile the model?

**Every time you make a change to your models, you need to recompile them.**

**In our case, we're using the exact same loss, optimizer and metrics as before, except this time the learning rate for our optimizer will be 10x smaller than before (0.0001 instead of Adam's default of 0.001).**

**We do this so the model doesn't try to overwrite the existing weights in the pretrained model too fast. In other words, we want learning to be more gradual.**

□ **Note:** There's no set standard for setting the learning rate during fine-tuning, though reductions of **2.6x-10x+ seem to work well in practice** .

How many trainable variables do we have now?

In [ ]:

```
print(len(model_2.trainable_variables))
```

12

Wonderful, it looks like our model has a total of 10 trainable variables, the last 10 layers of the base model and the weight and bias parameters of the Dense output layer.

Time to fine-tune!

We're going to continue training on from where our previous model finished. Since it trained for 5 epochs, our fine-tuning will begin on the epoch 5 and continue for another 5 epochs.

To do this, we can use the `initial_epoch` parameter of the `fit()` method. We'll pass it the last epoch of the previous model's training history (`history_10_percent_data_aug.epoch[-1]`).

In [ ]:

```
# Fine tune for another 5 epochs
fine_tune_epochs = initial_epochs + 5

# Refit the model (same as model_2 except with more trainable layers)
history_fine_10_percent_data_aug = model_2.fit(train_data_10_percent,
                                              epochs=fine_tune_epochs,
                                              validation_data=test_data,
                                              initial_epoch=history_10_percent_data_aug
                                              .epoch[-1], # start from previous last epoch
                                              validation_steps=int(0.25 * len(test_data
                                              )),
                                              callbacks=[create_tensorboard_callback("
transfer_learning", "10_percent_fine_tune_last_10")]) # name experiment appropriately
```

Saving TensorBoard log files to: transfer\_learning/10\_percent\_fine\_tune\_last\_10/20210216-022051

Epoch 5/10

24/24 [=====] - 15s 439ms/step - loss: 0.6963 - accuracy: 0.8062  
- val\_loss: 0.6032 - val\_accuracy: 0.8043

Epoch 6/10

24/24 [=====] - 8s 329ms/step - loss: 0.5747 - accuracy: 0.8390  
- val\_loss: 0.5580 - val\_accuracy: 0.8125

Epoch 7/10

24/24 [=====] - 8s 319ms/step - loss: 0.4972 - accuracy: 0.8458  
- val\_loss: 0.5543 - val\_accuracy: 0.8240

Epoch 8/10

24/24 [=====] - 8s 323ms/step - loss: 0.4262 - accuracy: 0.8814  
- val\_loss: 0.5403 - val\_accuracy: 0.8191

Epoch 9/10

24/24 [=====] - 8s 300ms/step - loss: 0.4234 - accuracy: 0.8855  
- val\_loss: 0.5262 - val\_accuracy: 0.8322

Epoch 10/10

24/24 [=====] - 8s 305ms/step - loss: 0.3665 - accuracy: 0.9056  
- val\_loss: 0.5218 - val\_accuracy: 0.8322

❏ **Note:** Fine-tuning usually takes far longer per epoch than feature extraction (due to updating more weights throughout a network).

Ho ho, looks like our model has gained a few percentage points of accuracy! Let's evaluate it.

In [ ]:

```
# Evaluate the model on the test data
results_fine_tune_10_percent = model_2.evaluate(test_data)
```

**Remember, the results from evaluating the model might be slightly different to the outputs from training since during training we only evaluate on 25% of the test data.**

**Alright, we need a way to evaluate our model's performance before and after fine-tuning. How about we write a function to compare the before and after?**

In [ ]:

```
def compare_historys(original_history, new_history, initial_epochs=5):
    """
    Compares two model history objects.
    """
    # Get original history measurements
    acc = original_history.history["accuracy"]
    loss = original_history.history["loss"]

    print(len(acc))

    val_acc = original_history.history["val_accuracy"]
    val_loss = original_history.history["val_loss"]

    # Combine original history with new history
    total_acc = acc + new_history.history["accuracy"]
    total_loss = loss + new_history.history["loss"]

    total_val_acc = val_acc + new_history.history["val_accuracy"]
    total_val_loss = val_loss + new_history.history["val_loss"]

    print(len(total_acc))
    print(total_acc)

    # Make plots
    plt.figure(figsize=(8, 8))
    plt.subplot(2, 1, 1)
    plt.plot(total_acc, label='Training Accuracy')
    plt.plot(total_val_acc, label='Validation Accuracy')
    plt.plot([initial_epochs-1, initial_epochs-1],
             plt.ylim(), label='Start Fine Tuning') # reshift plot around epochs
    plt.legend(loc='lower right')
    plt.title('Training and Validation Accuracy')


    plt.subplot(2, 1, 2)
    plt.plot(total_loss, label='Training Loss')
    plt.plot(total_val_loss, label='Validation Loss')
    plt.plot([initial_epochs-1, initial_epochs-1],
             plt.ylim(), label='Start Fine Tuning') # reshift plot around epochs
    plt.legend(loc='upper right')
    plt.title('Training and Validation Loss')
    plt.xlabel('epoch')
    plt.show()
```

**This is where saving the history variables of our model training comes in handy. Let's see what happened after fine-tuning the last 10 layers of our model.**

In [ ]:

```
compare_historys(original_history=history_10_percent_data_aug,
                 new_history=history_fine_10_percent_data_aug,
                 initial_epochs=5)
```

```
5
11
[0.32933333351612091, 0.64533333497047424, 0.73333333492279053, 0.76933333625793457, 0.794666
6479110718, 0.8026666641235352, 0.8320000171661377, 0.846666693687439, 0.8613333106040955
, 0.8893333077430725, 0.8893333077430725]
```



Alright, alright, seems like the curves are heading in the right direction after fine-tuning. But remember, it should be noted that fine-tuning usually works best with larger amounts of data.

## Model 4: Fine-tuning an existing model all of the data

Enough talk about how fine-tuning a model usually works with more data, let's try it out.

We'll start by downloading the full version of our 10 food classes dataset.

In [ ]:

```
# Download and unzip 10 classes of data with all images
!wget https://storage.googleapis.com/ztm_tf_course/food_vision/10_food_classes_all_data.zip
unzip_data("10_food_classes_all_data.zip")

# Setup data directories
train_dir = "10_food_classes_all_data/train/"
test_dir = "10_food_classes_all_data/test/"

--2021-02-16 02:48:30-- https://storage.googleapis.com/ztm_tf_course/food_vision/10_food_classes_all_data.zip
Resolving storage.googleapis.com (storage.googleapis.com)... 172.217.164.144, 172.253.115.128, 172.253.63.128, ...
Connecting to storage.googleapis.com (storage.googleapis.com)|172.217.164.144|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 519183241 (495M) [application/zip]
Saving to: '10_food_classes_all_data.zip'

10_food_classes_all 100%[=====>] 495.13M  124MB/s  in 4.1s

2021-02-16 02:48:35 (121 MB/s) - '10_food_classes_all_data.zip' saved [519183241/519183241]
```

In [ ]:

```
# How many images are we working with now?
walk_through_dir("10_food_classes_all_data")
```



```

There are 2 directories and 0 images in '10_food_classes_all_data'.
There are 10 directories and 0 images in '10_food_classes_all_data/train'.
There are 0 directories and 750 images in '10_food_classes_all_data/train/ice_cream'.
There are 0 directories and 750 images in '10_food_classes_all_data/train/ramen'.
There are 0 directories and 750 images in '10_food_classes_all_data/train/chicken_wings'.
There are 0 directories and 750 images in '10_food_classes_all_data/train/pizza'.
There are 0 directories and 750 images in '10_food_classes_all_data/train/steak'.
There are 0 directories and 750 images in '10_food_classes_all_data/train/fried_rice'.
There are 0 directories and 750 images in '10_food_classes_all_data/train/hamburger'.
There are 0 directories and 750 images in '10_food_classes_all_data/train/grilled_salmon'.
.
There are 0 directories and 750 images in '10_food_classes_all_data/train/sushi'.
There are 0 directories and 750 images in '10_food_classes_all_data/train/chicken_curry'.
There are 10 directories and 0 images in '10_food_classes_all_data/test'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/ice_cream'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/ramen'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/chicken_wings'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/pizza'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/steak'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/fried_rice'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/hamburger'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/grilled_salmon'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/sushi'.
There are 0 directories and 250 images in '10_food_classes_all_data/test/chicken_curry'.

```

**And now we'll turn the images into tensors datasets.**

In [ ]:

```

# Setup data inputs
import tensorflow as tf
IMG_SIZE = (224, 224)
train_data_10_classes_full = tf.keras.preprocessing.image_dataset_from_directory(train_dir,
                                                                                 label_
mode="categorical",
                                                                                 image_
size=IMG_SIZE)

# Note: this is the same test dataset we've been using for the previous modelling experim
ents
test_data = tf.keras.preprocessing.image_dataset_from_directory(test_dir,
                                                                label_mode="categorical
",
                                                                image_size=IMG_SIZE)

```

```

Found 7500 files belonging to 10 classes.
Found 2500 files belonging to 10 classes.

```

**Oh this is looking good. We've got 10x more images in of the training classes to work with.**

**The test dataset is the same we've been using for our previous experiments.**

**As it is now, our `model_2` has been fine-tuned on 10 percent of the data, so to begin fine-tuning on all of the data and keep our experiments consistent, we need to revert it back to the weights we checkpointed after 5 epochs of feature-extraction.**

**To demonstrate this, we'll first evaluate the current `model_2`.**

In [ ]:

```

# Evaluate model (this is the fine-tuned 10 percent of data version)
model_2.evaluate(test_data)

```

```

79/79 [=====] - 10s 118ms/step - loss: 0.4870 - accuracy: 0.8388

```

Out[ ]:

```

[0.4869591295719147, 0.8388000130653381]

```

These are the same values as `results_fine_tune_10_percent`.

In [ ]:

```
results_fine_tune_10_percent
```

Out[ ]:

```
[0.48695918917655945, 0.8388000130653381]
```

Now we'll revert the model back to the saved weights.

In [ ]:

```
# Load model from checkpoint, that way we can fine-tune from the same stage the 10 percent data model was fine-tuned from
model_2.load_weights(checkpoint_path) # revert model back to saved weights
```

Out[ ]:

```
<tensorflow.python.training.tracking.util.CheckpointLoadStatus at 0x7fcd9dale4a8>
```

And the results should be the same as `results_10_percent_data_aug`.

In [ ]:

```
# After loading the weights, this should have gone down (no fine-tuning)
model_2.evaluate(test_data)
```

```
79/79 [=====] - 10s 117ms/step - loss: 0.7047 - accuracy: 0.8080
```

Out[ ]:

```
[0.7046542763710022, 0.8080000281333923]
```

In [ ]:

```
# Check to see if the above two results are the same (they should be)
results_10_percent_data_aug
```

Out[ ]:

```
[0.7046541571617126, 0.8080000281333923]
```

Alright, the previous steps might seem quite confusing but all we've done is:

1. Trained a feature extraction transfer learning model for 5 epochs on 10% of the data (with all base model layers frozen) and saved the model's weights using `ModelCheckpoint`.
2. Fine-tuned the same model on the same 10% of the data for a further 5 epochs with the top 10 layers of the base model unfrozen.
3. Saved the results and training logs each time.
4. Reloaded the model from 1 to do the same steps as 2 but with all of the data.

The same steps as 2?

Yeah, we're going to fine-tune the last 10 layers of the base model with the full dataset for another 5 epochs but first let's remind ourselves which layers are trainable.

In [ ]:

```
# Check which layers are tuneable in the whole model
for layer_number, layer in enumerate(model_2.layers):
    print(layer_number, layer.name, layer.trainable)
```

```
0 input_layer True
1 data_augmentation True
2 efficientnetb0 True
3 global_average_pooling_layer True
4 output_layer True
```

## Can we get a little more specific?

In [ ]:

```
# Check which layers are tuneable in the base model
for layer_number, layer in enumerate(base_model.layers):
    print(layer_number, layer.name, layer.trainable)
```

```
0 input_3 False
1 rescaling_2 False
2 normalization_2 False
3 stem_conv_pad False
4 stem_conv False
5 stem_bn False
6 stem_activation False
7 block1a_dwconv False
8 block1a_bn False
9 block1a_activation False
10 block1a_se_squeeze False
11 block1a_se_reshape False
12 block1a_se_reduce False
13 block1a_se_expand False
14 block1a_se_excite False
15 block1a_project_conv False
16 block1a_project_bn False
17 block2a_expand_conv False
18 block2a_expand_bn False
19 block2a_expand_activation False
20 block2a_dwconv_pad False
21 block2a_dwconv False
22 block2a_bn False
23 block2a_activation False
24 block2a_se_squeeze False
25 block2a_se_reshape False
26 block2a_se_reduce False
27 block2a_se_expand False
28 block2a_se_excite False
29 block2a_project_conv False
30 block2a_project_bn False
31 block2b_expand_conv False
32 block2b_expand_bn False
33 block2b_expand_activation False
34 block2b_dwconv False
35 block2b_bn False
36 block2b_activation False
37 block2b_se_squeeze False
38 block2b_se_reshape False
39 block2b_se_reduce False
40 block2b_se_expand False
41 block2b_se_excite False
42 block2b_project_conv False
43 block2b_project_bn False
44 block2b_drop False
45 block2b_add False
46 block3a_expand_conv False
47 block3a_expand_bn False
48 block3a_expand_activation False
49 block3a_dwconv_pad False
50 block3a_dwconv False
51 block3a_bn False
52 block3a_activation False
53 block3a_se_squeeze False
54 block3a_se_reshape False
55 block3a_se_reduce False
56 block3a_se_expand False
57 block3a_se_excite False
58 block3a_project_conv False
59 block3a_project_bn False
60 block3b_expand_conv False
61 block3b_expand_bn False
62 block3b_expand_activation False
```

63 block3b\_dwconv False  
64 block3b\_bn False  
65 block3b\_activation False  
66 block3b\_se\_squeeze False  
67 block3b\_se\_reshape False  
68 block3b\_se\_reduce False  
69 block3b\_se\_expand False  
70 block3b\_se\_excite False  
71 block3b\_project\_conv False  
72 block3b\_project\_bn False  
73 block3b\_drop False  
74 block3b\_add False  
75 block4a\_expand\_conv False  
76 block4a\_expand\_bn False  
77 block4a\_expand\_activation False  
78 block4a\_dwconv\_pad False  
79 block4a\_dwconv False  
80 block4a\_bn False  
81 block4a\_activation False  
82 block4a\_se\_squeeze False  
83 block4a\_se\_reshape False  
84 block4a\_se\_reduce False  
85 block4a\_se\_expand False  
86 block4a\_se\_excite False  
87 block4a\_project\_conv False  
88 block4a\_project\_bn False  
89 block4b\_expand\_conv False  
90 block4b\_expand\_bn False  
91 block4b\_expand\_activation False  
92 block4b\_dwconv False  
93 block4b\_bn False  
94 block4b\_activation False  
95 block4b\_se\_squeeze False  
96 block4b\_se\_reshape False  
97 block4b\_se\_reduce False  
98 block4b\_se\_expand False  
99 block4b\_se\_excite False  
100 block4b\_project\_conv False  
101 block4b\_project\_bn False  
102 block4b\_drop False  
103 block4b\_add False  
104 block4c\_expand\_conv False  
105 block4c\_expand\_bn False  
106 block4c\_expand\_activation False  
107 block4c\_dwconv False  
108 block4c\_bn False  
109 block4c\_activation False  
110 block4c\_se\_squeeze False  
111 block4c\_se\_reshape False  
112 block4c\_se\_reduce False  
113 block4c\_se\_expand False  
114 block4c\_se\_excite False  
115 block4c\_project\_conv False  
116 block4c\_project\_bn False  
117 block4c\_drop False  
118 block4c\_add False  
119 block5a\_expand\_conv False  
120 block5a\_expand\_bn False  
121 block5a\_expand\_activation False  
122 block5a\_dwconv False  
123 block5a\_bn False  
124 block5a\_activation False  
125 block5a\_se\_squeeze False  
126 block5a\_se\_reshape False  
127 block5a\_se\_reduce False  
128 block5a\_se\_expand False  
129 block5a\_se\_excite False  
130 block5a\_project\_conv False  
131 block5a\_project\_bn False  
132 block5b\_expand\_conv False  
133 block5b\_expand\_bn False  
134 block5b\_expand\_activation False

135 block5b\_dwconv False  
136 block5b\_bn False  
137 block5b\_activation False  
138 block5b\_se\_squeeze False  
139 block5b\_se\_reshape False  
140 block5b\_se\_reduce False  
141 block5b\_se\_expand False  
142 block5b\_se\_excite False  
143 block5b\_project\_conv False  
144 block5b\_project\_bn False  
145 block5b\_drop False  
146 block5b\_add False  
147 block5c\_expand\_conv False  
148 block5c\_expand\_bn False  
149 block5c\_expand\_activation False  
150 block5c\_dwconv False  
151 block5c\_bn False  
152 block5c\_activation False  
153 block5c\_se\_squeeze False  
154 block5c\_se\_reshape False  
155 block5c\_se\_reduce False  
156 block5c\_se\_expand False  
157 block5c\_se\_excite False  
158 block5c\_project\_conv False  
159 block5c\_project\_bn False  
160 block5c\_drop False  
161 block5c\_add False  
162 block6a\_expand\_conv False  
163 block6a\_expand\_bn False  
164 block6a\_expand\_activation False  
165 block6a\_dwconv\_pad False  
166 block6a\_dwconv False  
167 block6a\_bn False  
168 block6a\_activation False  
169 block6a\_se\_squeeze False  
170 block6a\_se\_reshape False  
171 block6a\_se\_reduce False  
172 block6a\_se\_expand False  
173 block6a\_se\_excite False  
174 block6a\_project\_conv False  
175 block6a\_project\_bn False  
176 block6b\_expand\_conv False  
177 block6b\_expand\_bn False  
178 block6b\_expand\_activation False  
179 block6b\_dwconv False  
180 block6b\_bn False  
181 block6b\_activation False  
182 block6b\_se\_squeeze False  
183 block6b\_se\_reshape False  
184 block6b\_se\_reduce False  
185 block6b\_se\_expand False  
186 block6b\_se\_excite False  
187 block6b\_project\_conv False  
188 block6b\_project\_bn False  
189 block6b\_drop False  
190 block6b\_add False  
191 block6c\_expand\_conv False  
192 block6c\_expand\_bn False  
193 block6c\_expand\_activation False  
194 block6c\_dwconv False  
195 block6c\_bn False  
196 block6c\_activation False  
197 block6c\_se\_squeeze False  
198 block6c\_se\_reshape False  
199 block6c\_se\_reduce False  
200 block6c\_se\_expand False  
201 block6c\_se\_excite False  
202 block6c\_project\_conv False  
203 block6c\_project\_bn False  
204 block6c\_drop False  
205 block6c\_add False  
206 block6d\_expand\_conv False

```

207 block6d_expand_bn False
208 block6d_expand_activation False
209 block6d_dwconv False
210 block6d_bn False
211 block6d_activation False
212 block6d_se_squeeze False
213 block6d_se_reshape False
214 block6d_se_reduce False
215 block6d_se_expand False
216 block6d_se_excite False
217 block6d_project_conv False
218 block6d_project_bn False
219 block6d_drop False
220 block6d_add False
221 block7a_expand_conv False
222 block7a_expand_bn False
223 block7a_expand_activation False
224 block7a_dwconv False
225 block7a_bn False
226 block7a_activation False
227 block7a_se_squeeze True
228 block7a_se_reshape True
229 block7a_se_reduce True
230 block7a_se_expand True
231 block7a_se_excite True
232 block7a_project_conv True
233 block7a_project_bn True
234 top_conv True
235 top_bn True
236 top_activation True

```

**Looking good! The last 10 layers are trainable (unfrozen).**

**We've got one more step to do before we can begin fine-tuning.**

**Do you remember what it is?**

**I'll give you a hint. We just reloaded the weights to our model and what do we need to do every time we make a change to our models?**

**Recompile them!**

**This will be just as before.**

In [ ]:

```

# Compile
model_2.compile(loss="categorical_crossentropy",
                 optimizer=tf.keras.optimizers.Adam(lr=0.0001), # divide learning rate by
10 for fine-tuning
                 metrics=["accuracy"])

```

**Alright, time to fine-tune on all of the data!**

In [ ]:

```

# Continue to train and fine-tune the model to our data
fine_tune_epochs = initial_epochs + 5

history_fine_10_classes_full = model_2.fit(train_data_10_classes_full,
                                           epochs=fine_tune_epochs,
                                           initial_epoch=history_10_percent_data_aug.epochs[-1],
                                           validation_data=test_data,
                                           validation_steps=int(0.25 * len(test_data)),
                                           callbacks=[create_tensorboard_callback("transfer_learning", "full_10_classes_fine_tune_last_10")])

```

Saving TensorBoard log files to: transfer\_learning/full\_10\_classes\_fine\_tune\_last\_10/20210216-025031

Transfer Learning

```

Epoch 5/10
235/235 [=====] - 49s 190ms/step - loss: 0.7943 - accuracy: 0.75
13 - val_loss: 0.4029 - val_accuracy: 0.8635
Epoch 6/10
235/235 [=====] - 51s 215ms/step - loss: 0.6391 - accuracy: 0.79
13 - val_loss: 0.3668 - val_accuracy: 0.8882
Epoch 7/10
235/235 [=====] - 47s 198ms/step - loss: 0.5564 - accuracy: 0.82
20 - val_loss: 0.3237 - val_accuracy: 0.9013
Epoch 8/10
235/235 [=====] - 46s 192ms/step - loss: 0.5030 - accuracy: 0.83
94 - val_loss: 0.3390 - val_accuracy: 0.8832
Epoch 9/10
235/235 [=====] - 45s 191ms/step - loss: 0.4611 - accuracy: 0.84
79 - val_loss: 0.3099 - val_accuracy: 0.8980
Epoch 10/10
235/235 [=====] - 43s 183ms/step - loss: 0.4507 - accuracy: 0.85
57 - val_loss: 0.2903 - val_accuracy: 0.9161

```

❏ **Note:** Training took longer per epoch, but that makes sense because we're using 10x more training data than before.

Let's evaluate on all of the test data.

In [ ]:

```

results_fine_tune_full_data = model_2.evaluate(test_data)
results_fine_tune_full_data

```

```

79/79 [=====] - 10s 118ms/step - loss: 0.3264 - accuracy: 0.8988

```

Out[ ]:

```

[0.32638612389564514, 0.8988000154495239]

```

Nice! It looks like fine-tuning with all of the data has given our model a boost, how do the training curves look?

In [ ]:

```


# How did fine-tuning go with more data?
compare_historys(original_history=history_10_percent_data_aug,
                  new_history=history_fine_10_classes_full,
                  initial_epochs=5)


```

```

5
11
[0.3293333351612091, 0.6453333497047424, 0.7333333492279053, 0.7693333625793457, 0.794666
6479110718, 0.7675999999046326, 0.8066666722297668, 0.8333333134651184, 0.847866654396057
1, 0.8511999845504761, 0.8655999898910522]

```





Looks like that extra data helped! Those curves are looking great. And if we trained for longer, they might even keep improving.

## Viewing our experiment data on TensorBoard

Right now our experimental results are scattered all throughout our notebook. If we want to share them with someone, they'd be getting a bunch of different graphs and metrics... not a fun time.

But guess what?

Thanks to the TensorBoard callback we made with our helper function `create_tensorflow_callback()`, we've been tracking our modelling experiments the whole time.

How about we upload them to TensorBoard.dev and check them out?

We can do with the `tensorboard dev upload` command and passing it the directory where our experiments have been logged.

**Note:** Remember, whatever you upload to TensorBoard.dev becomes public. If there are training logs you don't want to share, don't upload them.

In [ ]:

```
# View tensorboard logs of transfer learning modelling experiments (should be 4 models)
# Upload TensorBoard dev records
!tensorboard dev upload --logdir ./transfer_learning \
  --name "Transfer learning experiments" \
  --description "A series of different transfer learning experiments with varying amounts
of data and fine-tuning" \
  --one_shot # exits the uploader when upload has finished
```

```
2020-09-17 22:51:36.043126: I tensorflow/stream_executor/platform/default/dso_loader.cc:4
8] Successfully opened dynamic library libcudart.so.10.1
Data for the "graphs" plugin is now uploaded to TensorBoard.dev! Note that uploaded data
is public. If you do not want to upload data for this plugin, use the "--plugins" command
line argument.
Data for the "histograms" plugin is now uploaded to TensorBoard.dev! Note that uploaded d
ata is public. If you do not want to upload data for this plugin, use the "--plugins" comm
and line argument.
Data for the "hparams" plugin is now uploaded to TensorBoard.dev! Note that uploaded data
is public. If you do not want to upload data for this plugin, use the "--plugins" command
line argument.
Upload started and will continue reading any new data as it's added
to the logdir. To stop uploading, press Ctrl-C.
```

View your TensorBoard live at: <https://tensorboard.dev/experiment/2076kw3PQbKl0lByfg5B4w/>

```
[2020-09-17T22:51:37] Uploader started.
[2020-09-17T22:51:47] Total uploaded: 128 scalars, 0 tensors, 5 binary objects (9.1 MB)
Listening for new data in logdir...
Done. View your TensorBoard at https://tensorboard.dev/experiment/2076kw3PQbKl0lByfg5B4w/
```

Once we've uploaded the results to TensorBoard.dev we get a shareable link we can use to view and compare our experiments and share our results with others if needed.



You can view the original versions of the experiments we ran in this notebook here:

<https://tensorboard.dev/experiment/2O76kw3PQbKl0lByfg5B4w/>

**Question: Which model performed the best? Why do you think this is? How did fine-tuning go?**

To find all of your previous TensorBoard.dev experiments using the command `tensorboard dev list`.

In [ ]:

```
# View previous experiments
```

```
!tensorboard dev list
```

```
2020-09-17 22:51:48.747476: I tensorflow/stream_executor/platform/default/dso_loader.cc:4
8] Successfully opened dynamic library libcudart.so.10.1
Data for the "graphs" plugin is now uploaded to TensorBoard.dev! Note that uploaded data
is public. If you do not want to upload data for this plugin, use the "--plugins" command
line argument.
Data for the "histograms" plugin is now uploaded to TensorBoard.dev! Note that uploaded d
ata is public. If you do not want to upload data for this plugin, use the "--plugins" comm
and line argument.
Data for the "hparams" plugin is now uploaded to TensorBoard.dev! Note that uploaded data
is public. If you do not want to upload data for this plugin, use the "--plugins" command
line argument.
https://tensorboard.dev/experiment/2O76kw3PQbKl0lByfg5B4w/
Name                Transfer learning experiments
Description          A series of different transfer learning experiments with varying am
ounts of data and fine-tuning
Id                  2O76kw3PQbKl0lByfg5B4w
Created             2020-09-17 22:51:37 (15 seconds ago)
Updated            2020-09-17 22:51:47 (5 seconds ago)
Runs               10
Tags               3
Scalars            128
Tensor bytes       0
Binary object bytes 9520961
https://tensorboard.dev/experiment/73taSKxXQeGPQsNBcVvY3g/
Name                EfficientNetB0 vs. ResNet50V2
Description          Comparing two different TF Hub feature extraction models architectu
res using 10% of training images
Id                  73taSKxXQeGPQsNBcVvY3g
Created             2020-09-14 05:02:48
Updated            2020-09-14 05:02:50
Runs               4
Tags               3
Scalars            40
Tensor bytes       0
Binary object bytes 3402042
Total: 2 experiment(s)
```

And if you want to remove a previous experiment (and delete it from public viewing) you can use the command:

```
tensorboard dev delete --experiment_id [INSERT_EXPERIMENT_ID_TO_DELETE]
```

In [ ]:

```
# Remove previous experiments
```

```
#!tensorboard dev delete --experiment_id OUbW003pRqqQgAphVBxi8Q
```

```
2020-09-17 22:51:53.982454: I tensorflow/stream_executor/platform/default/dso_loader.cc:4
8] Successfully opened dynamic library libcudart.so.10.1
Data for the "graphs" plugin is now uploaded to TensorBoard.dev! Note that uploaded data
is public. If you do not want to upload data for this plugin, use the "--plugins" command
line argument.
Data for the "histograms" plugin is now uploaded to TensorBoard.dev! Note that uploaded d
ata is public. If you do not want to upload data for this plugin, use the "--plugins" comm
and line argument.
Data for the "hparams" plugin is now uploaded to TensorBoard.dev! Note that uploaded data
is public. If you do not want to upload data for this plugin, use the "--plugins" command
line argument.
```

No such experiment OUbW003pRqqQgAphVBxi8Q. Either it never existed or it has already been deleted.

## ▮ Exercises

1. Write a function to visualize an image from any dataset (train or test file) and any class (e.g. "steak", "pizza"... etc), visualize it and make a prediction on it using a trained model.
2. Use feature-extraction to train a transfer learning model on 10% of the Food Vision data for 10 epochs using `tf.keras.applications.EfficientNetB0` as the base model. Use the `ModelCheckpoint` callback to save the weights to file.
3. Fine-tune the last 20 layers of the base model you trained in 2 for another 10 epochs. How did it go?
4. Fine-tune the last 30 layers of the base model you trained in 2 for another 10 epochs. How did it go?

## ▮ Extra-curriculum

- Read the [documentation on data augmentation](#) in TensorFlow.
- Read the [ULMFit paper](#) (technical) for an introduction to the concept of freezing and unfreezing different layers.
- Read up on learning rate scheduling (there's a [TensorFlow callback](#) for this), how could this influence our model training?
  - If you're training for longer, you probably want to reduce the learning rate as you go... the closer you get to the bottom of the hill, the smaller steps you want to take. Imagine it like finding a coin at the bottom of your couch. In the beginning your arm movements are going to be large and the closer you get, the smaller your movements become.